
CIS 471/571 (Fall 2020): Introduction to Artificial Intelligence

Assignment Project Exam Help

Lecture 6: Adversarial Search

<https://tutorcs.com>

WeChat: cstutorcs

Thanh H. Nguyen

Source: <http://ai.berkeley.edu/home.html>



Reminders

- Project 2:
 - Deadline: Oct 27th, 2020
 - Written assignment 2:
 - Deadline: Oct 24th, 2020
- Assignment Project Exam Help
- <https://tutorcs.com>
- WeChat: cstutorcs

Adversarial Games



Assignment Project Exam Help

<https://tutores.com>

WeChat: cstutores



Types of Games

- Many different kinds of games!
- Axes:
 - Deterministic or stochastic?
 - One, two, or more players?
 - Zero sum?
 - Perfect information (can you see the state)?
- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Deterministic Games

- Many possible formalizations, one is:
 - States: S (start at s_0)
 - Players: $P=\{1...N\}$ (usually take turns)
 - Actions: A (may depend on player / state)
 - Transition Function: $S \times A \rightarrow S$
 - Terminal Test: $S \rightarrow \{t, f\}$
 - Terminal Utilities: $S \times P \rightarrow R$
- Solution for a player is a **policy**: $S \rightarrow A$


Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Zero-Sum Games



■ Zero-Sum Games

- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition

■ General Games


- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible
- More later on non-zero-sum games

Adversarial Search


Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

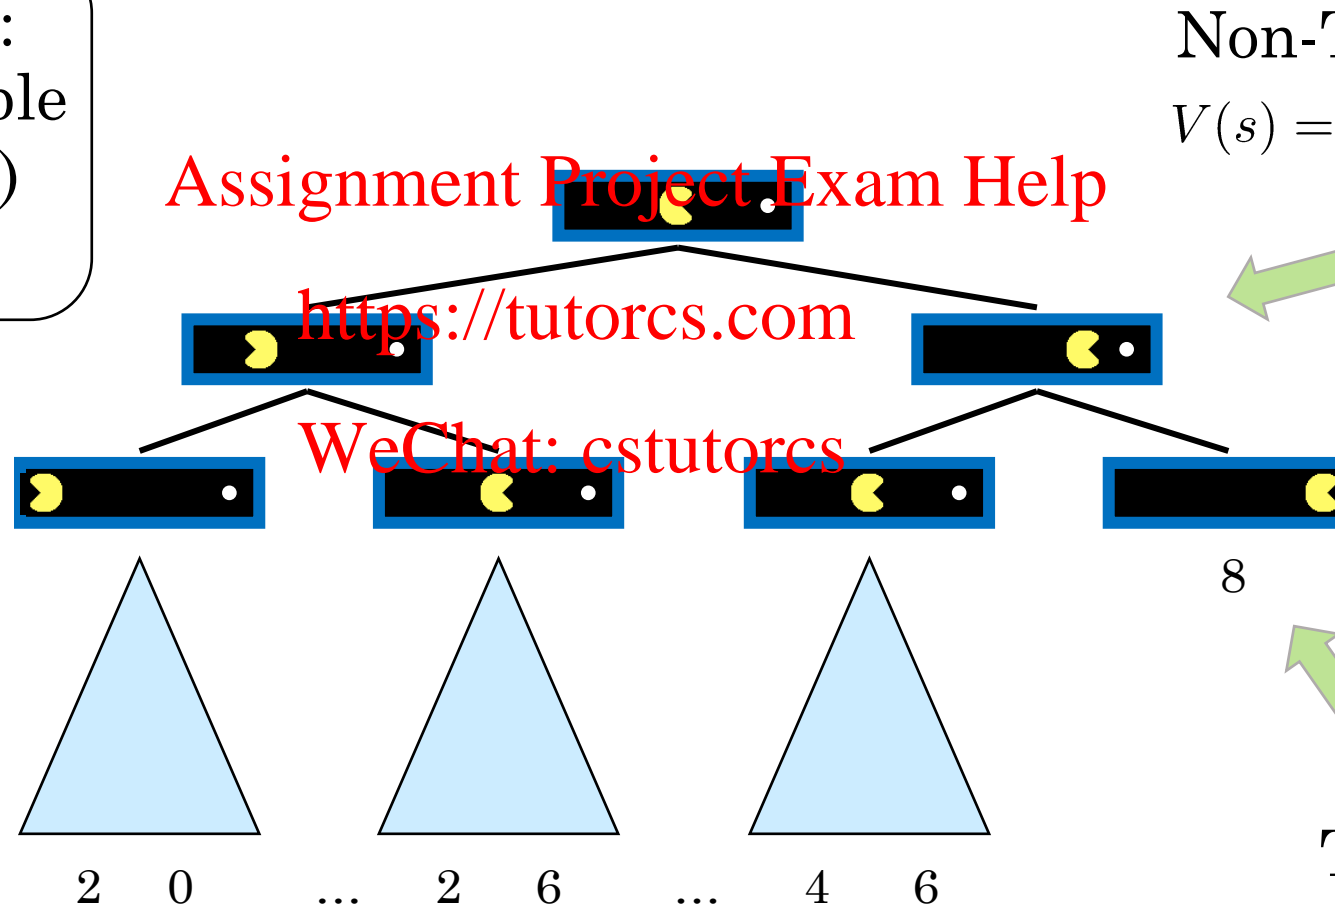


Single-Agent Trees



Value of a State

Value of a state:
The best achievable
outcome (utility)
from that state



Non-Terminal States:

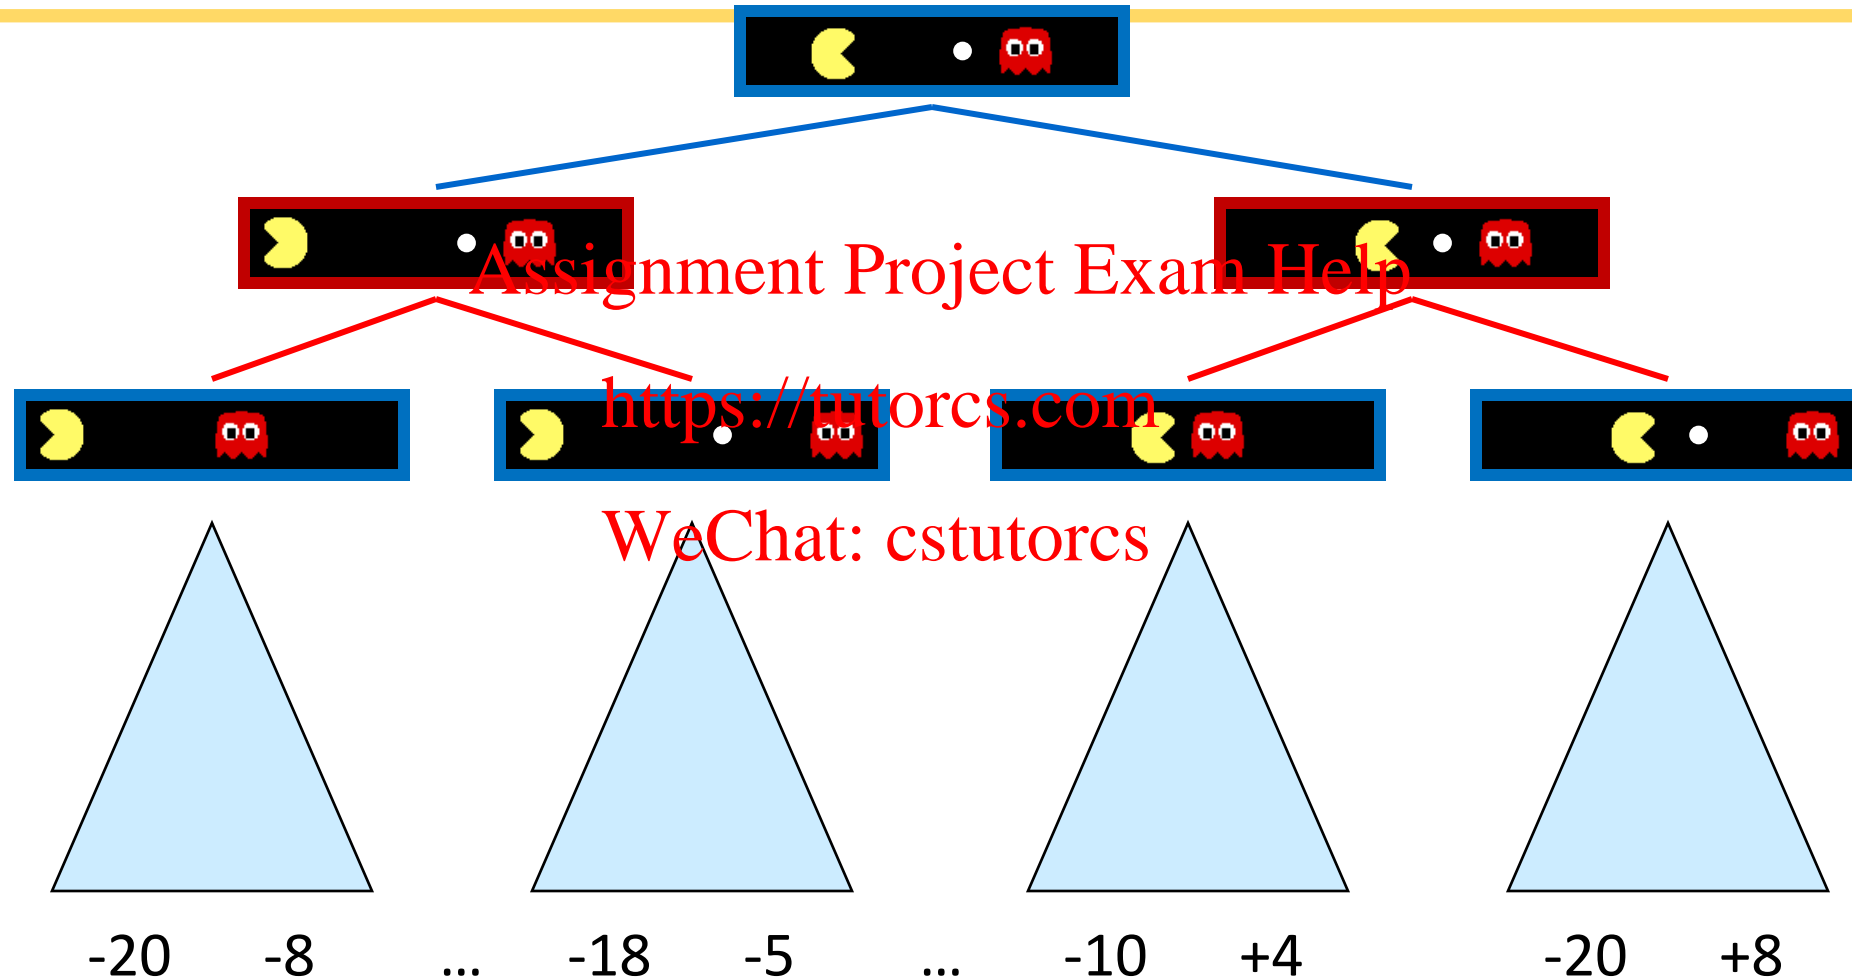
$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

Terminal States:

$$V(s) = \text{known}$$



Adversarial Game Trees



Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$


States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Terminal States:

$$V(s) = \text{known}$$



Adversarial Search (Minimax)

- Deterministic, zero-sum games:

- Tic-tac-toe, chess, checkers
- One player maximizes result
- The other minimizes result

Assignment Project Exam Help


<https://tutorcs.com>

- Minimax search:

- A state-space search tree
- Players alternate turns
- Compute each node's **minimax value**:
the best achievable utility against a
rational (optimal) adversary

WeChat: cstutorcs

Minimax values:
computed recursively



Terminal values:
part of the game



Minimax Implementation

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

<https://tutorcs.com>

WeChat: cstutorcs

def max-value(state):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v


def min-value(state):

initialize $v = +\infty$

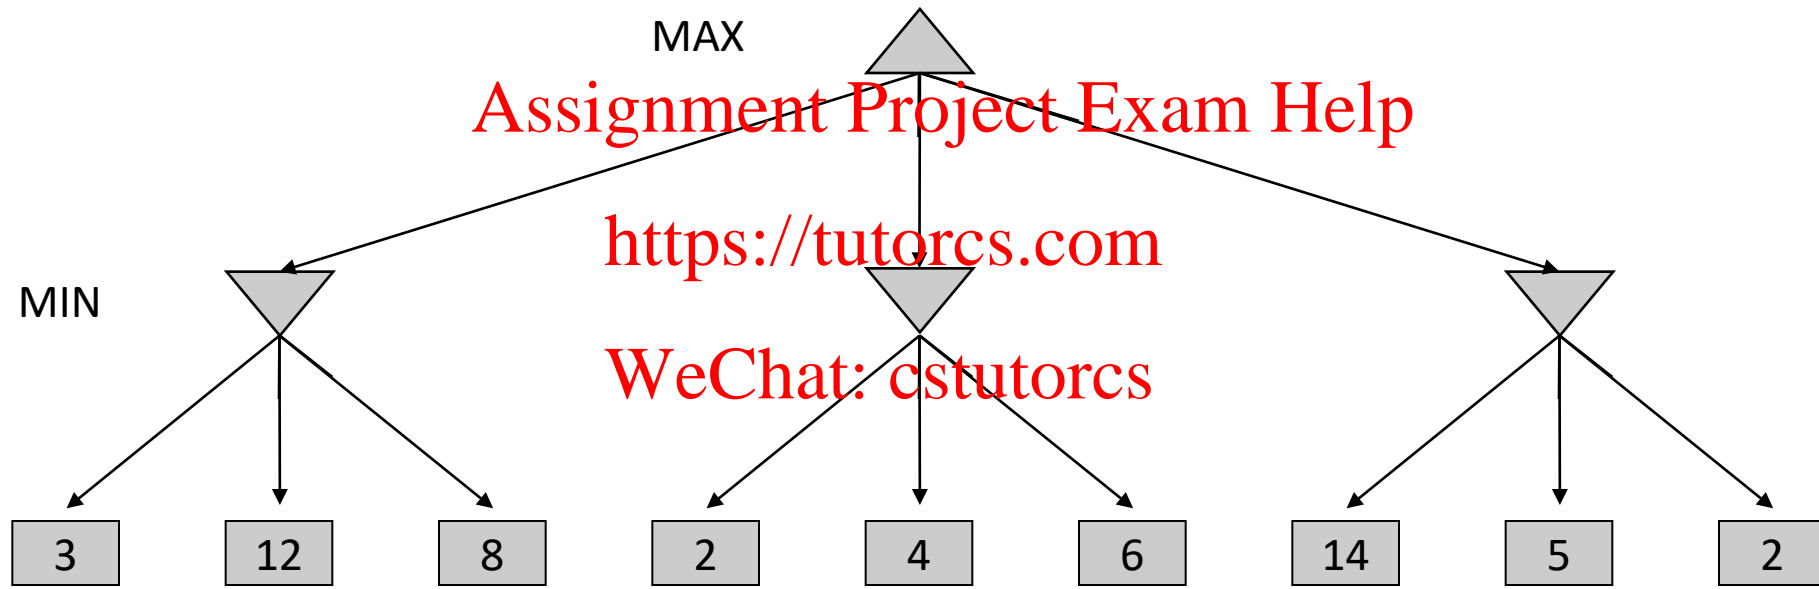
for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$


return v



Minimax Example





Minimax Properties



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Optimal against a perfect player. Otherwise?



Minimax Efficiency

- How efficient is minimax?

- Just like (exhaustive) DFS
- Time: $O(b^m)$
- Space: $O(bm)$


Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- Example: For chess, $b \approx 35$, $m \approx 100$

- Exact solution is completely infeasible
- But, do we need to explore the whole tree?



Resource Limits


Assignment Project Exam Help

<https://tutorcs.com>

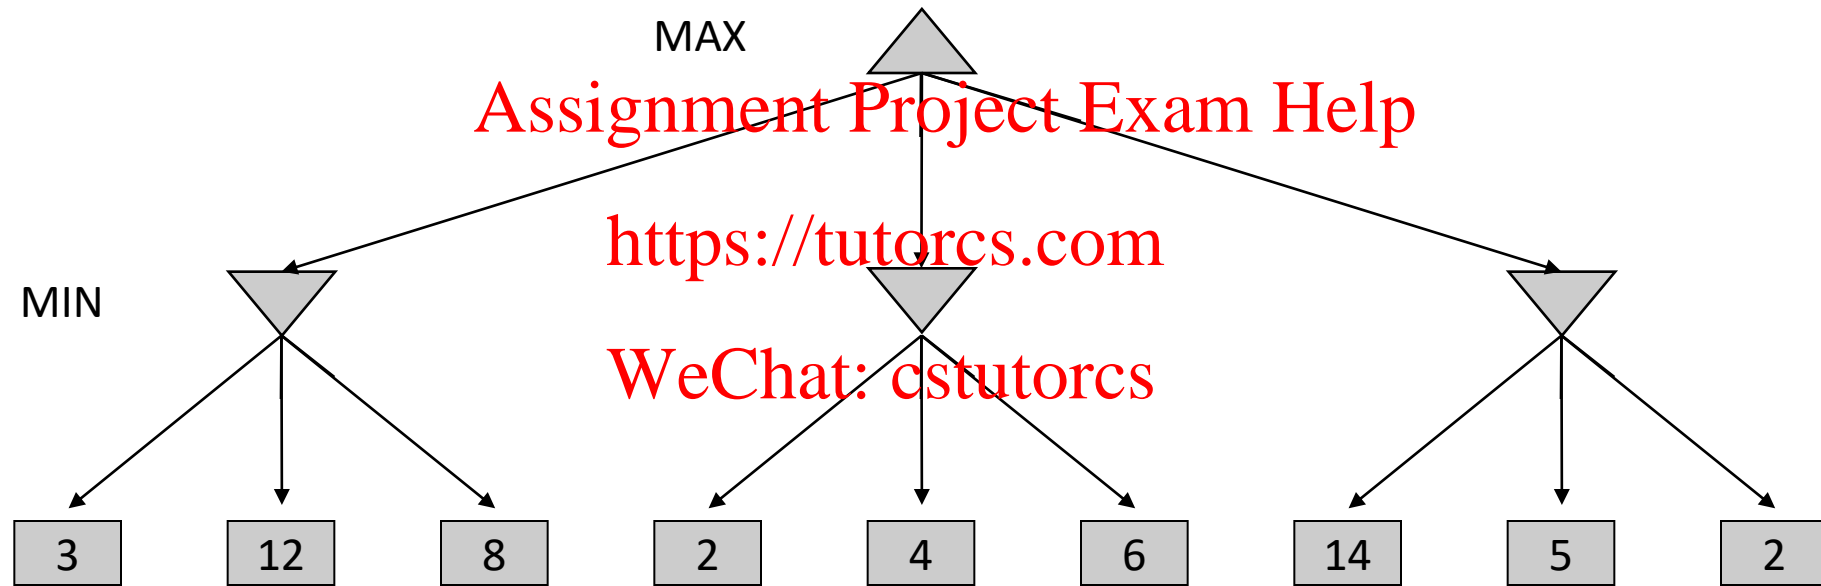
WeChat: cstutorcs



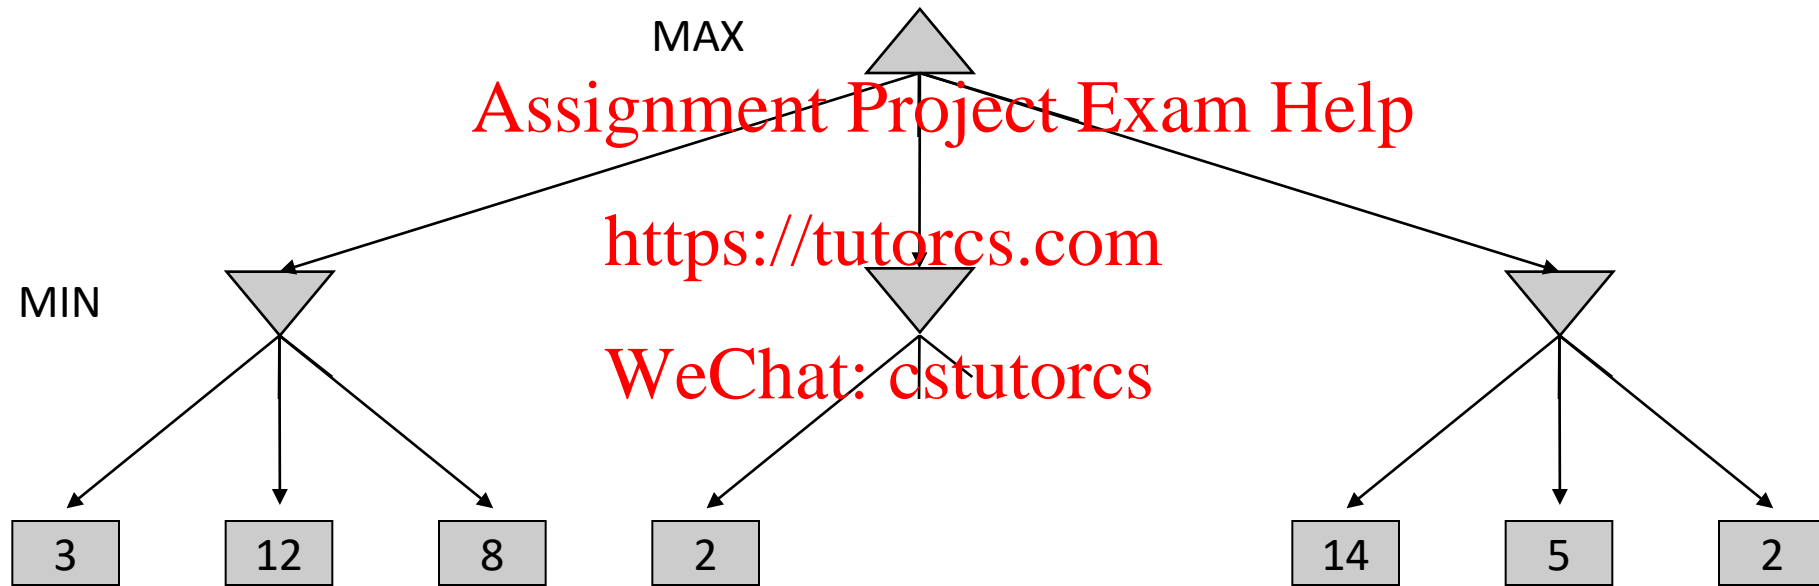
Game Tree Pruning



Minimax Example



Minimax Pruning



Alpha-Beta Pruning

- **Alpha α** : value of the best choice so far for MAX (lower bound of Max utility)
- **Beta β** : value of the best choice so far for MIN (upper bound of Min utility)
- Expanding at MAX node **n**: update α
 - If a child of **n** has value greater than β , stop expanding the MAX node **n**
 - Reason: MIN parent of **n** would not choose the action which leads to **n**
- At MIN node **n**: update β
 - If a child of **n** has value less than α , stop expanding the MIN node **n**
 - Reason: MAX parent of **n** would not choose the action which leads to **n**

WeChat: cstutorcs

Alpha-Beta Implementation

def value(state, α , β):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state, α , β)

if the next agent is MIN: return min-value(state, α , β)

<https://tutorcs.com>

def max-value(state, α , β):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \geq \beta$ return v

$\alpha = \max(\alpha, v)$

return v

WeChat: cstutorcs

def min-value(state, α , β):

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


if $v \leq \alpha$ return v

$\beta = \min(\beta, v)$

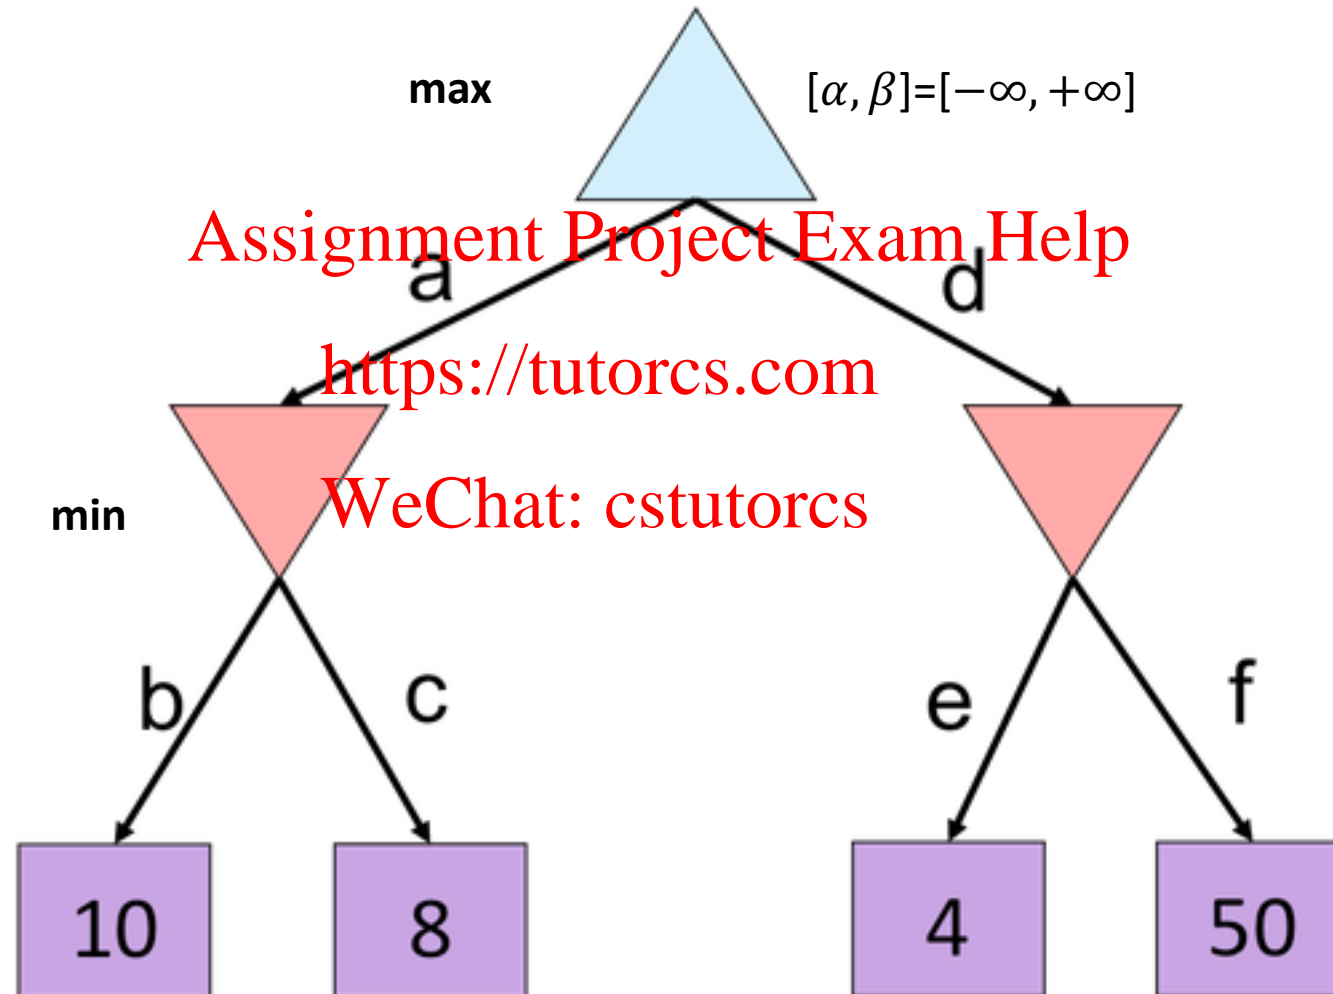
return v

Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - So the most naïve version won't let you do action selection
- Good child ordering improves effectiveness of pruning



Alpha-Beta Quiz



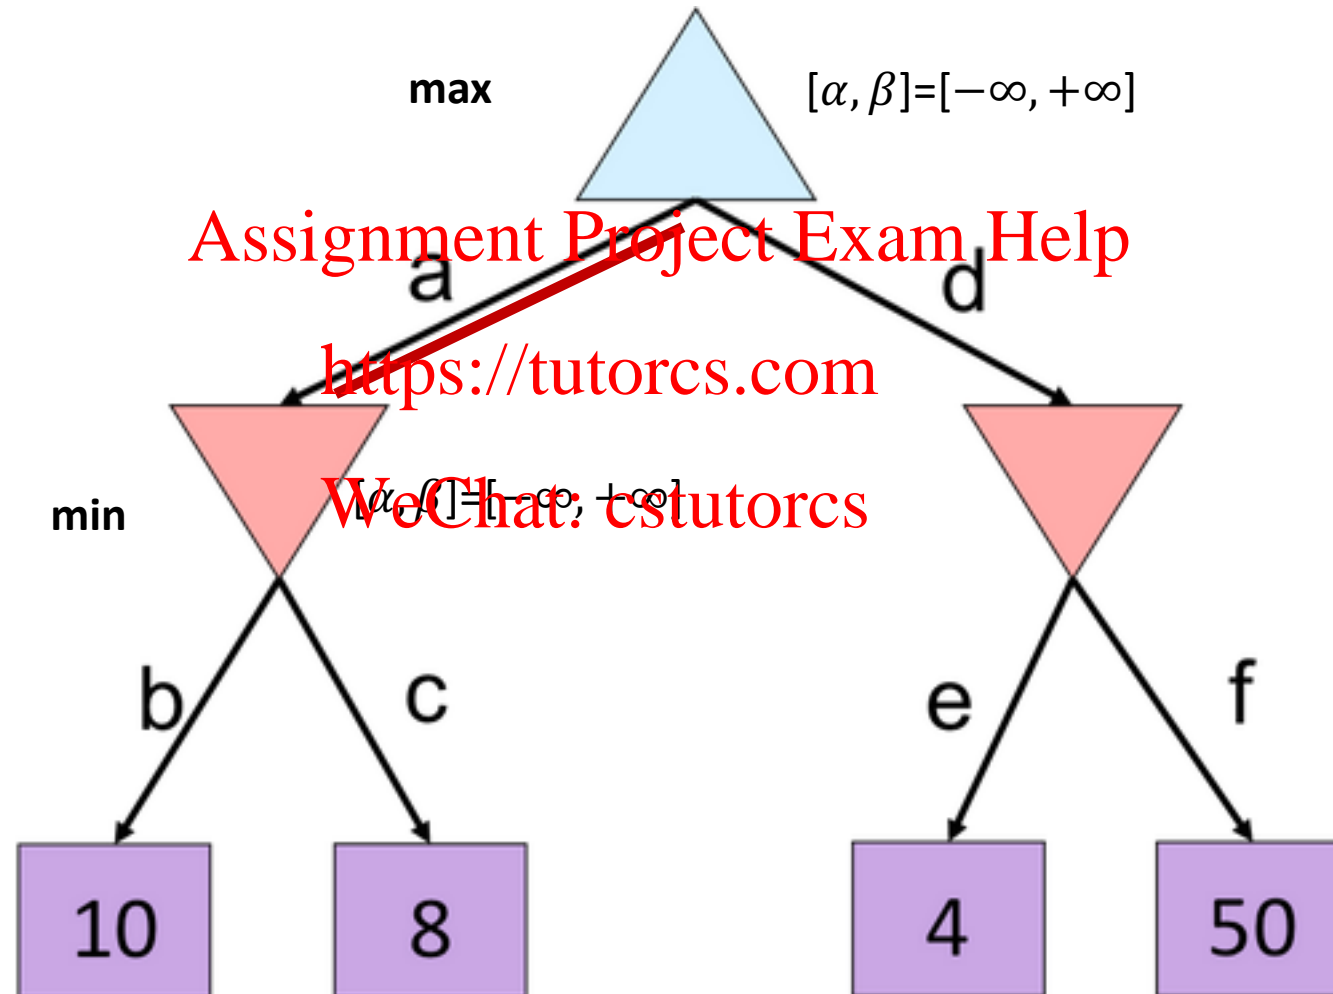
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Alpha-Beta Quiz



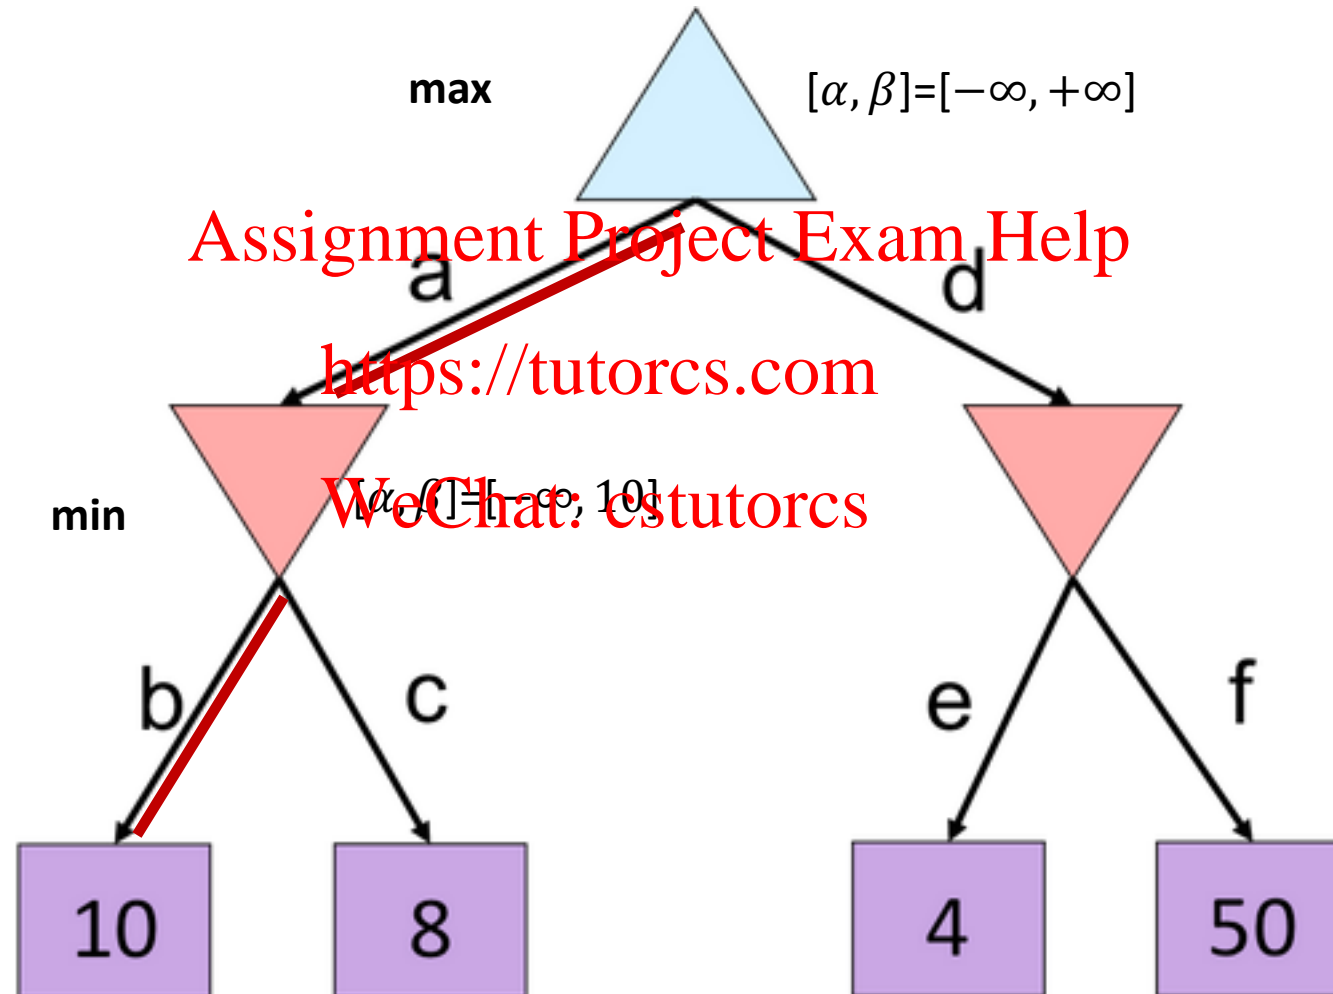
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Alpha-Beta Quiz



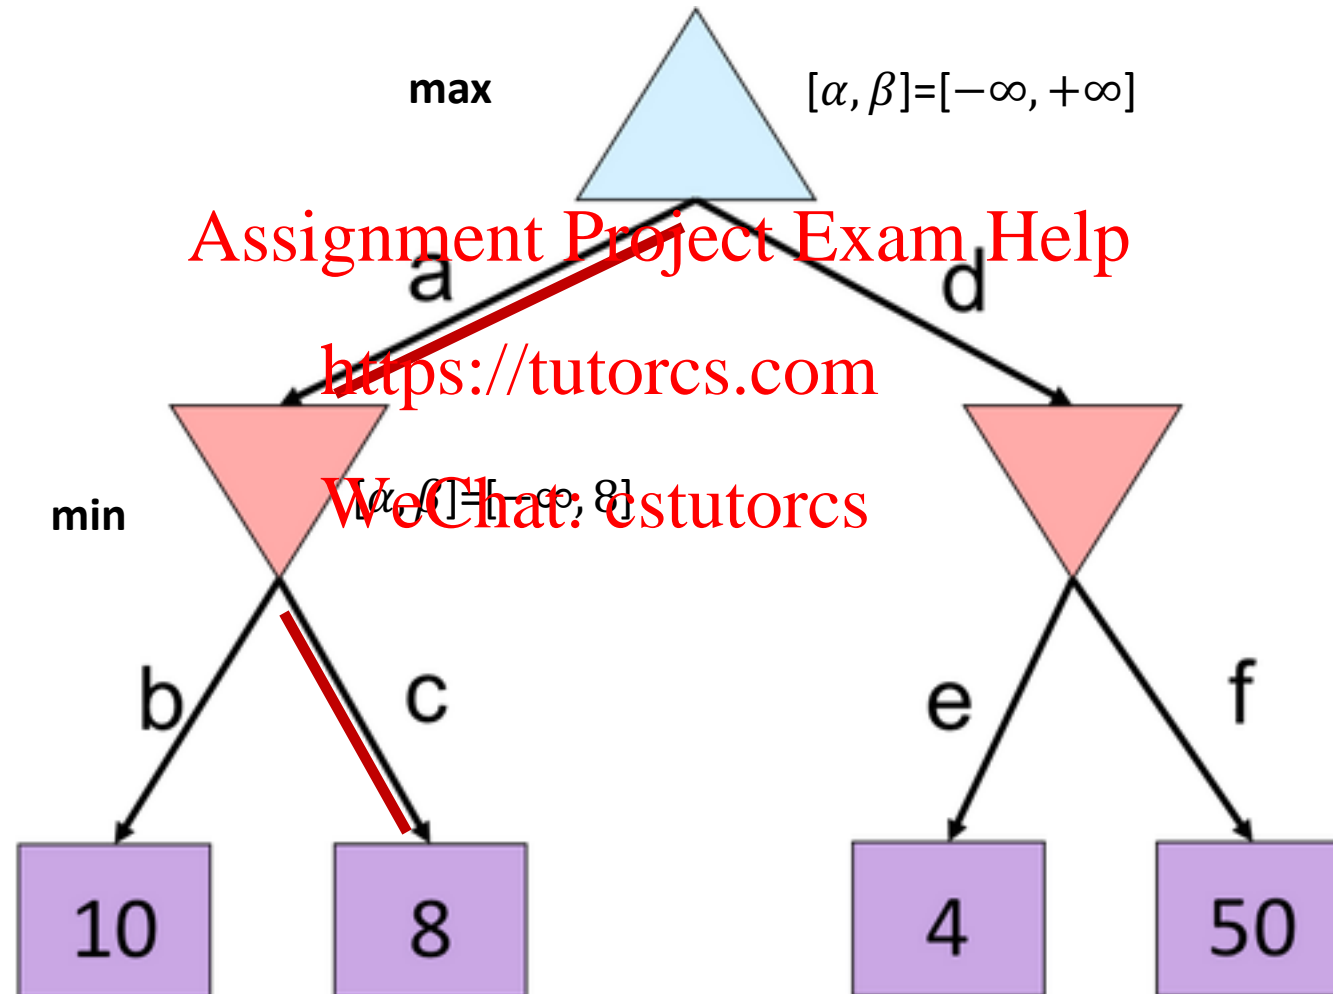
Assignment Project Exam Help

<https://tutorcs.com>

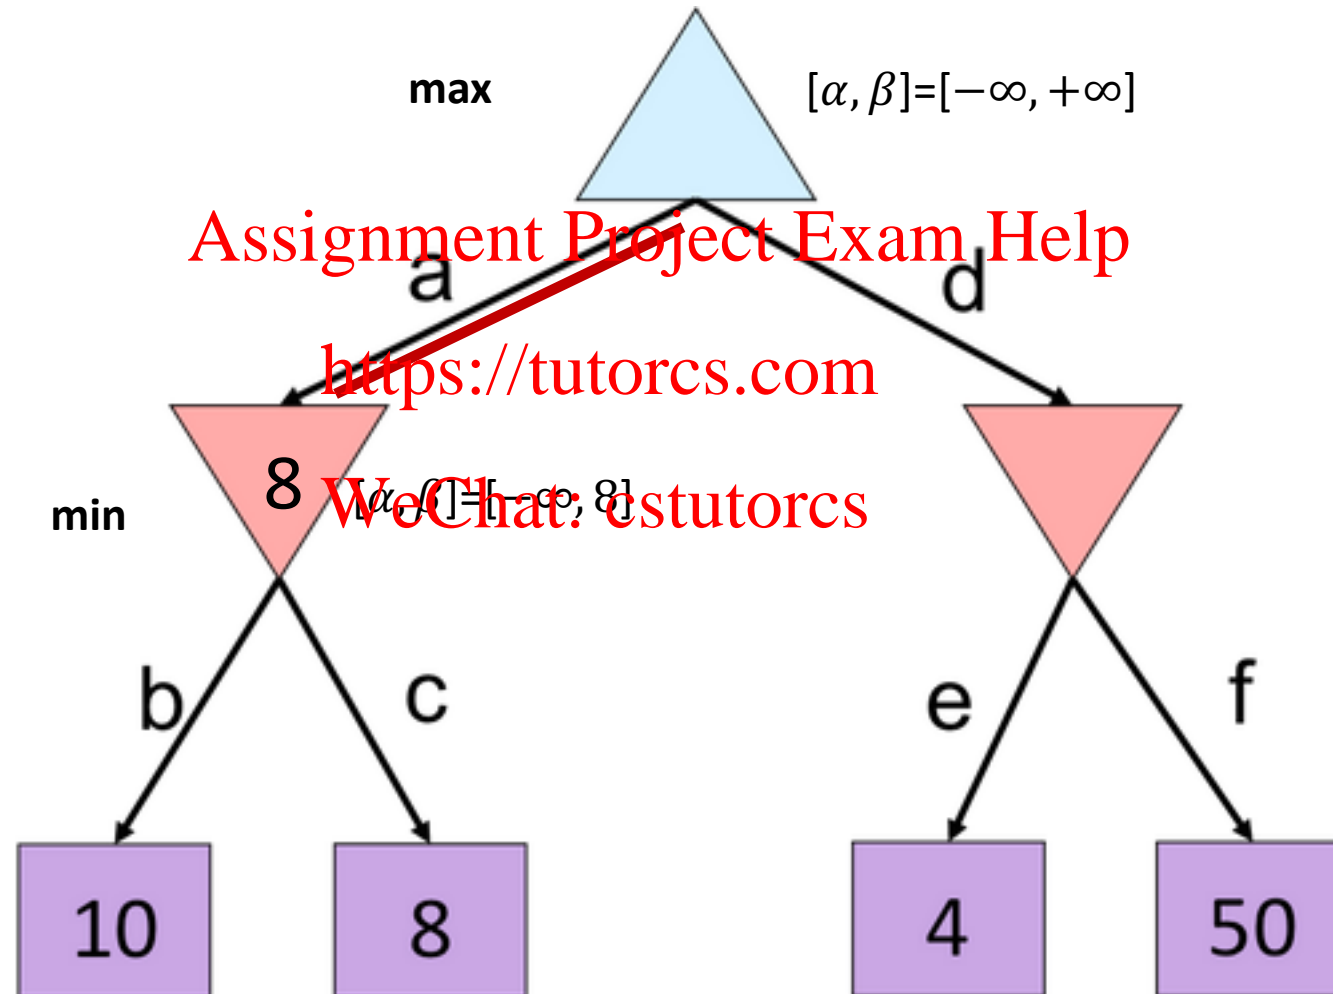
WeChat: estutorcs



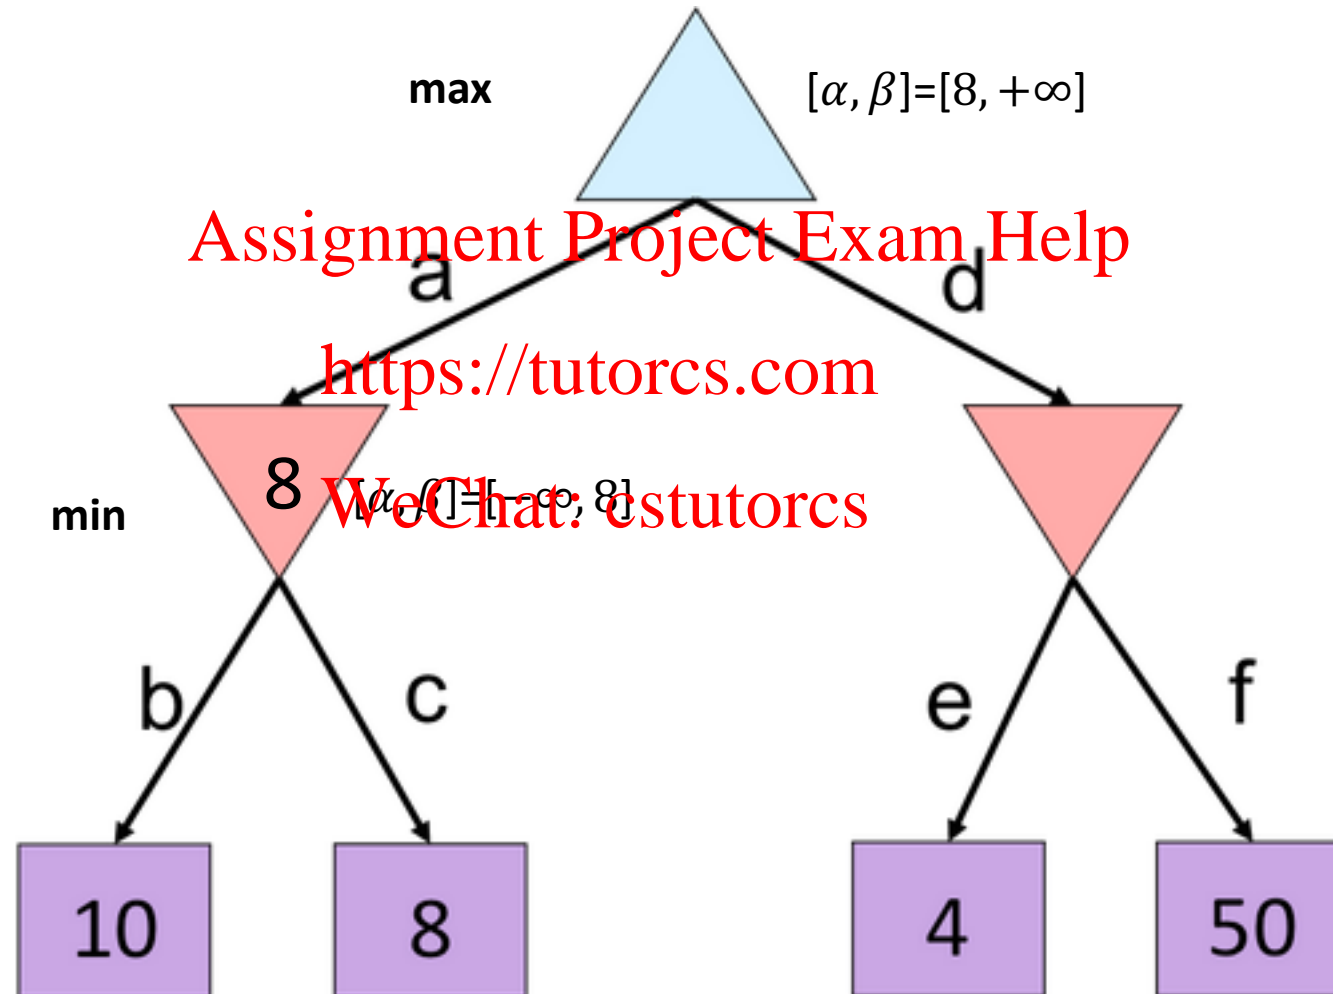
Alpha-Beta Quiz



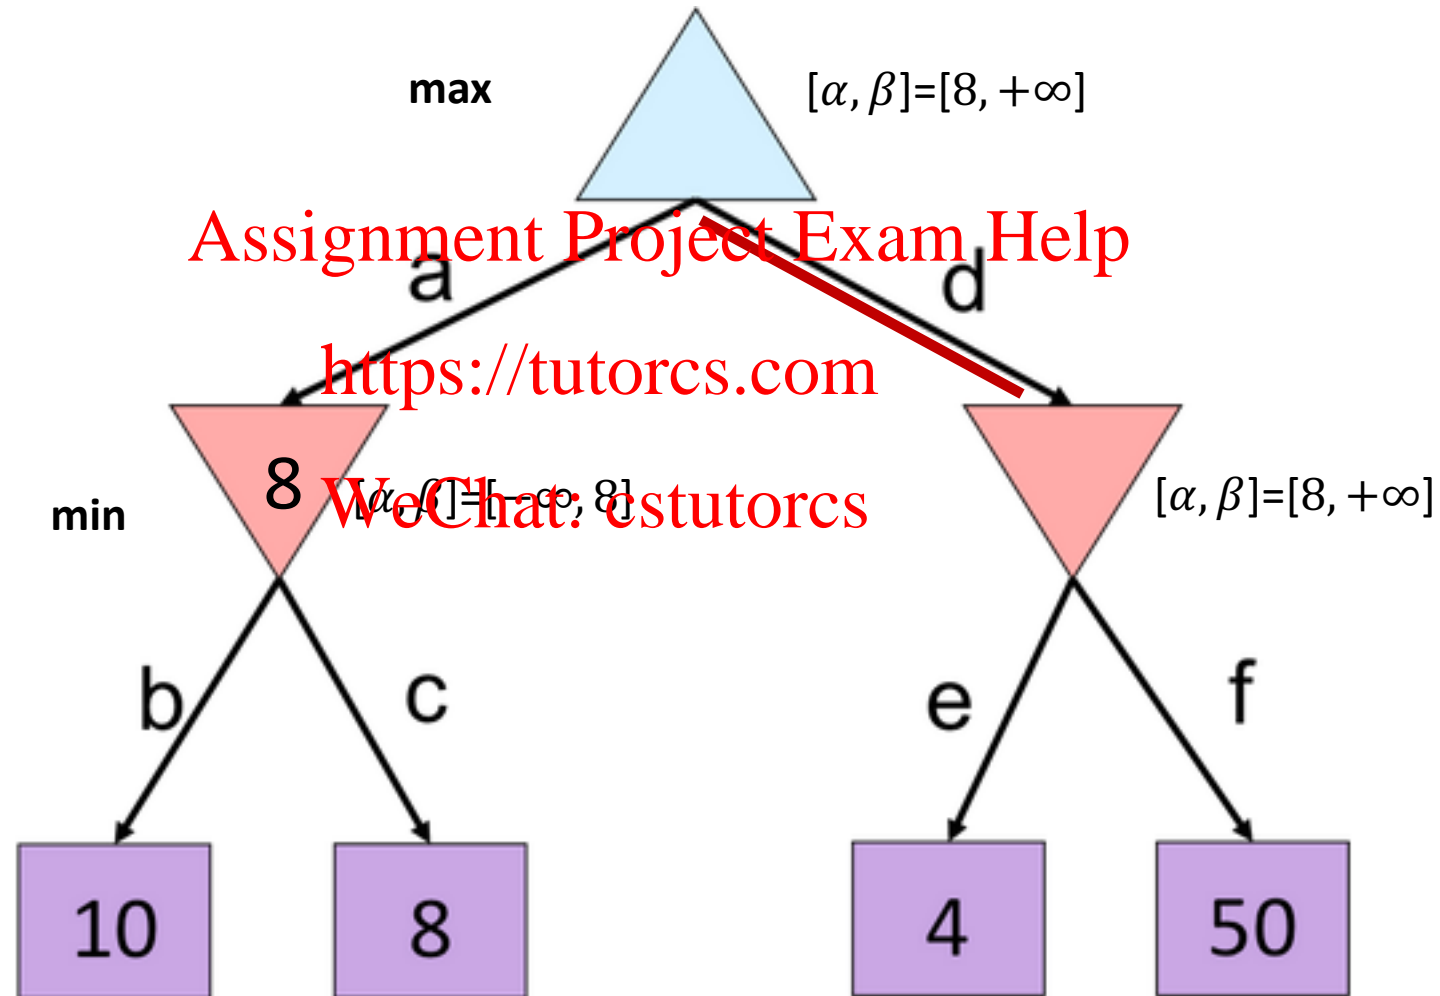
Alpha-Beta Quiz



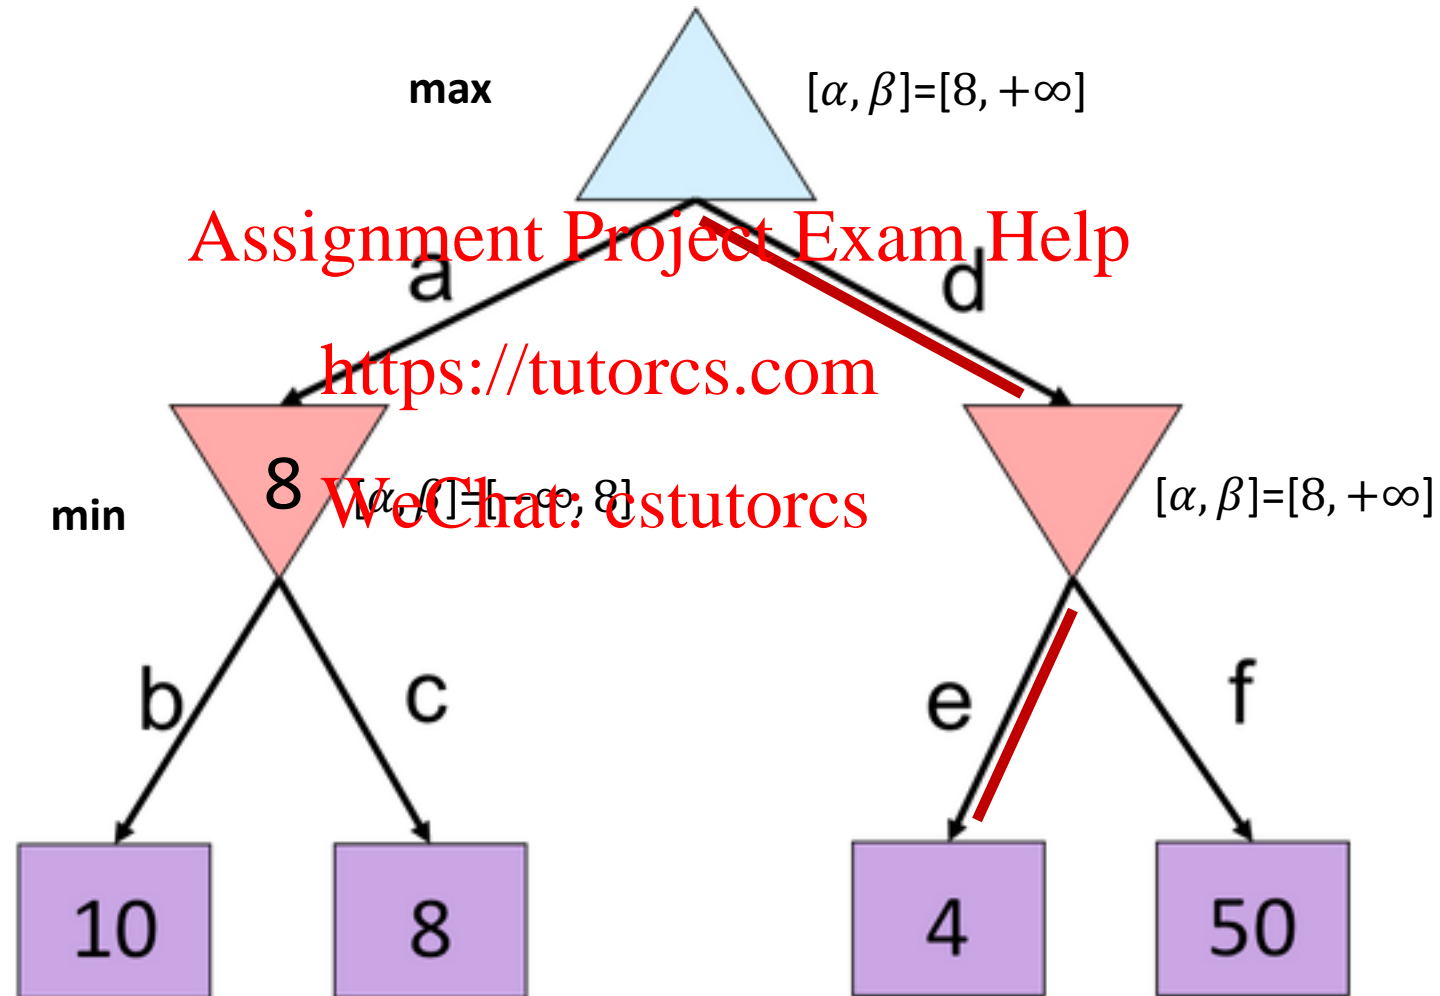
Alpha-Beta Quiz



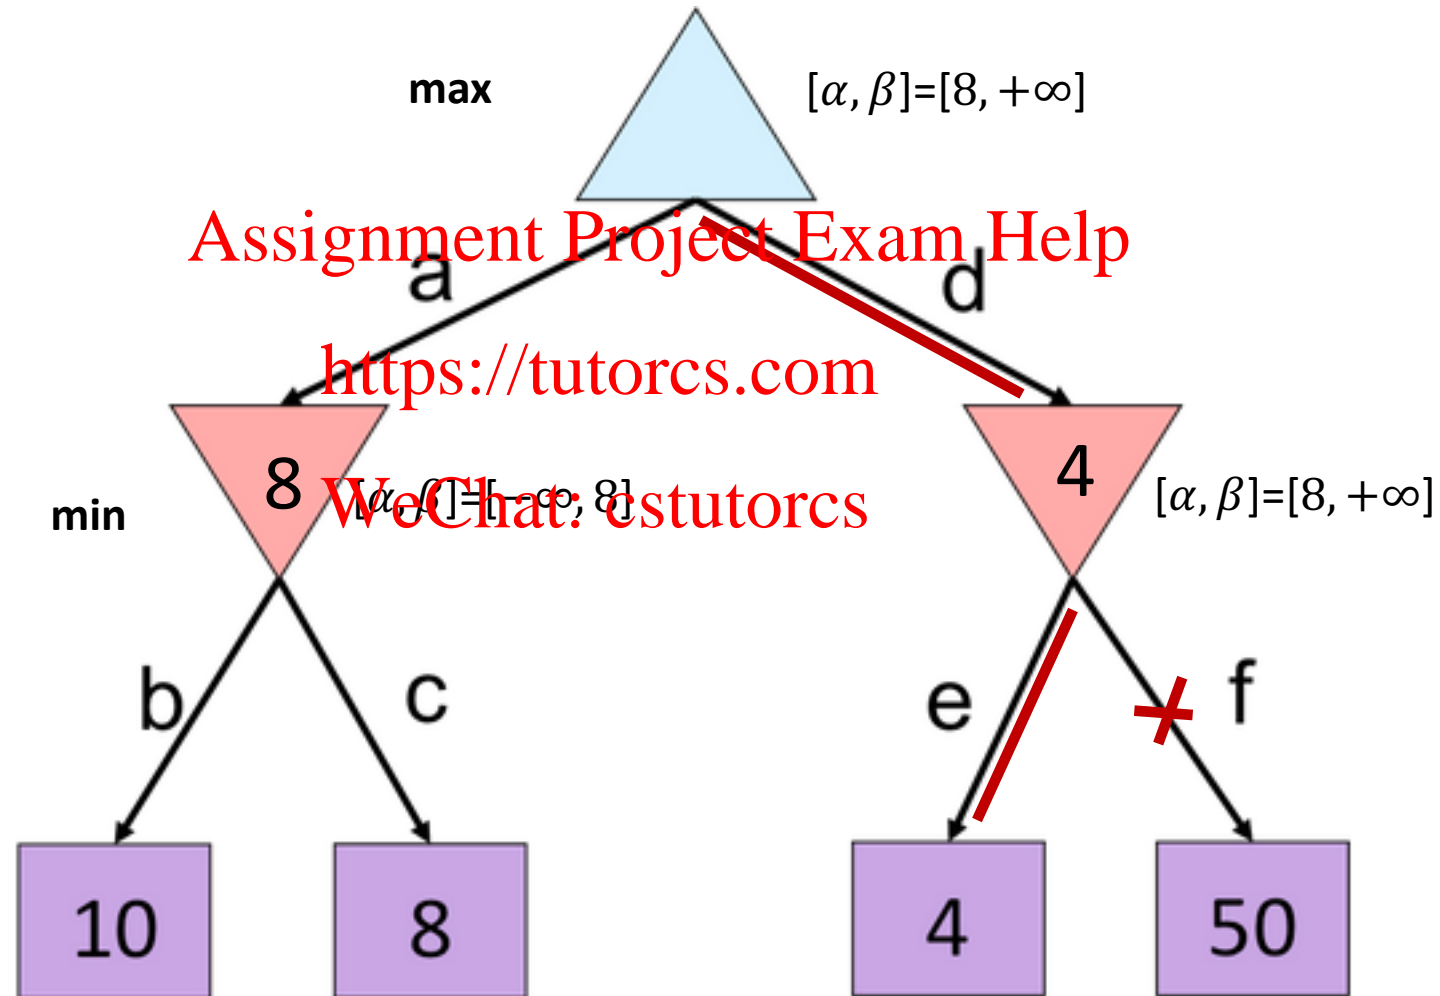
Alpha-Beta Quiz



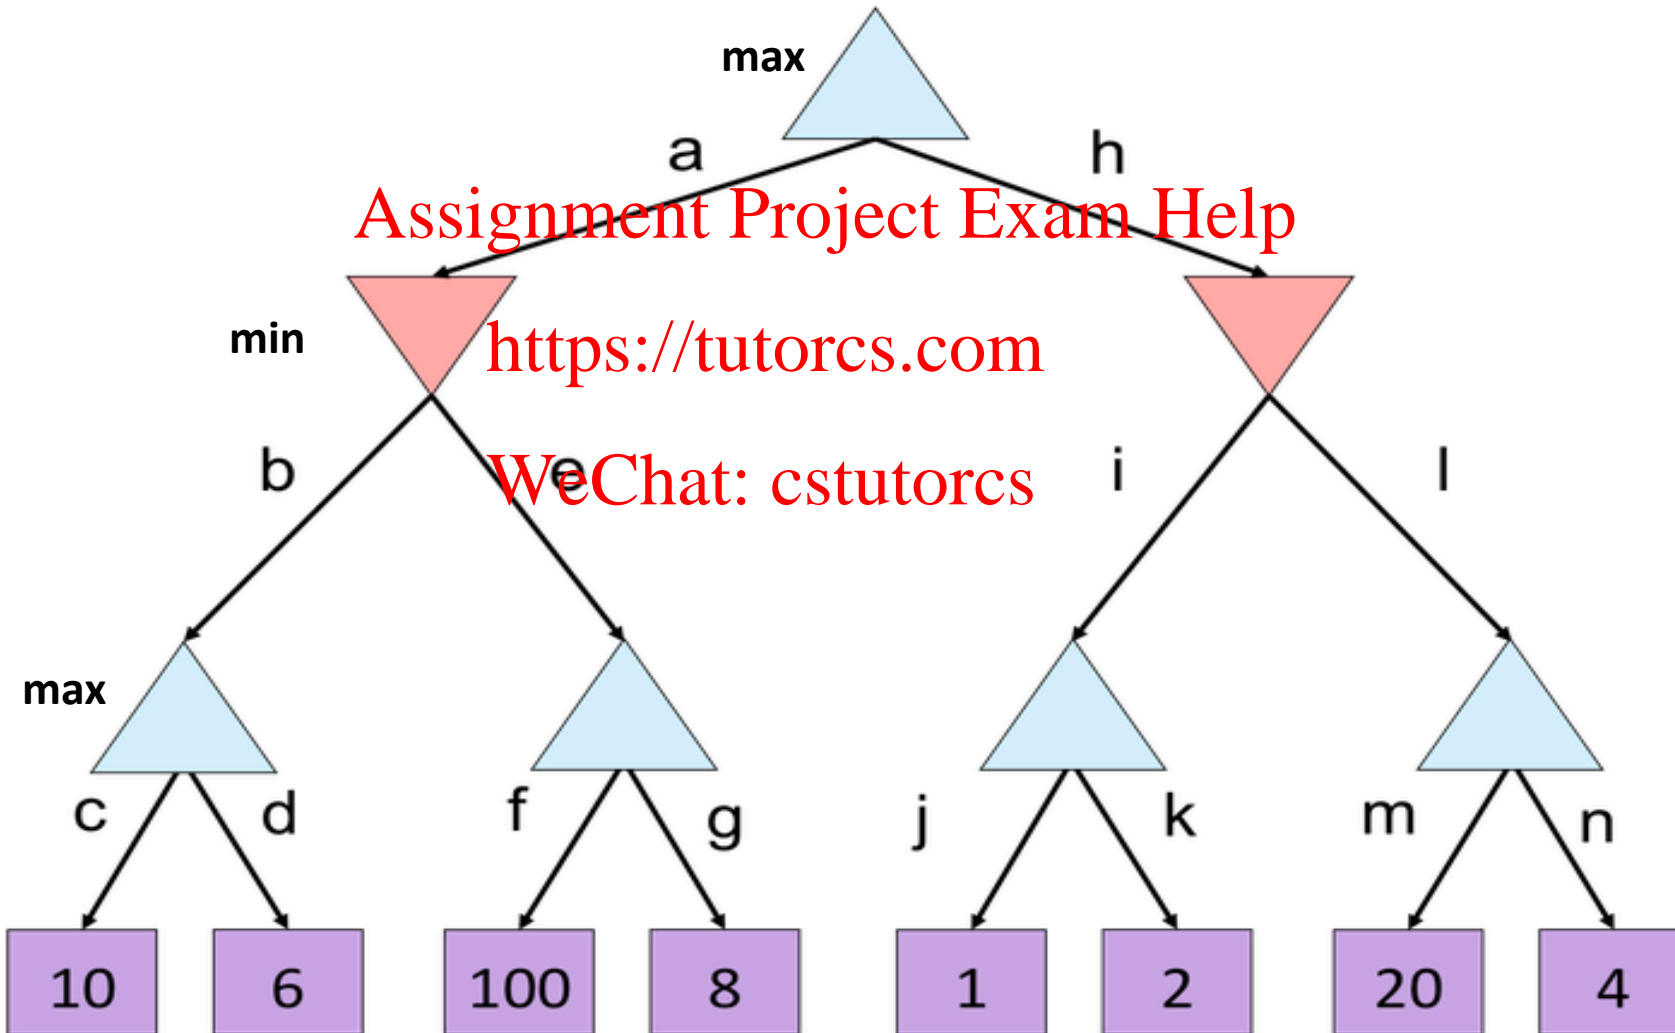
Alpha-Beta Quiz



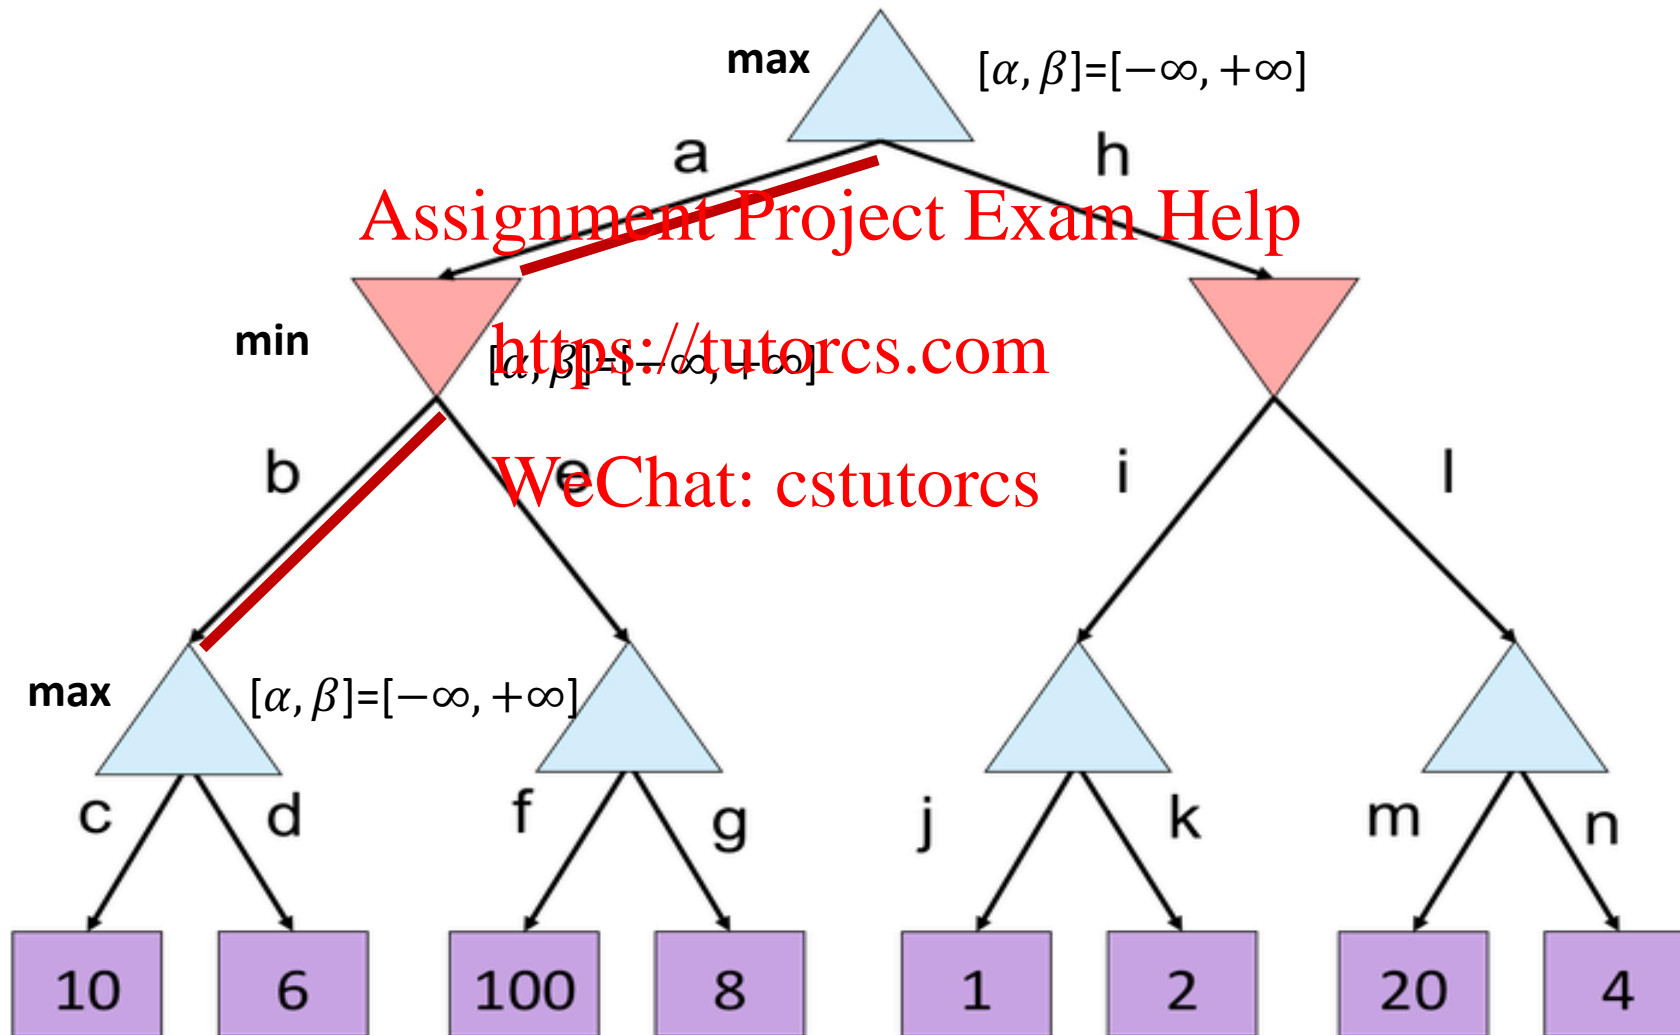
Alpha-Beta Quiz



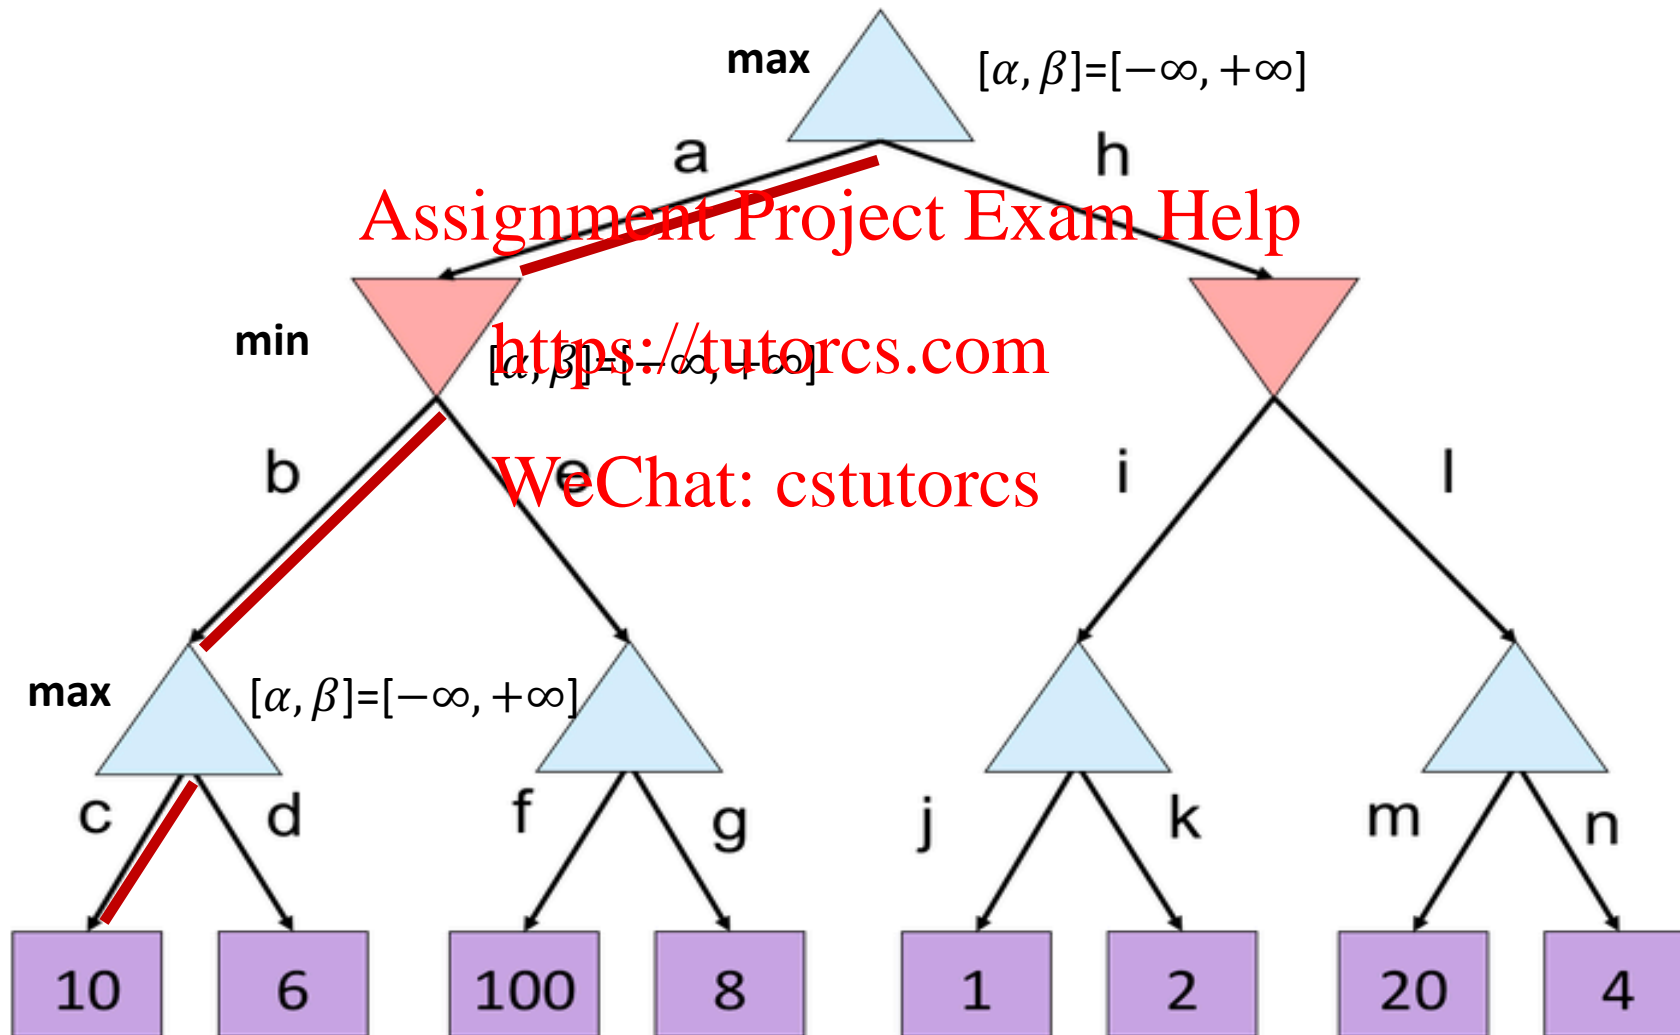
Alpha-Beta Quiz 2



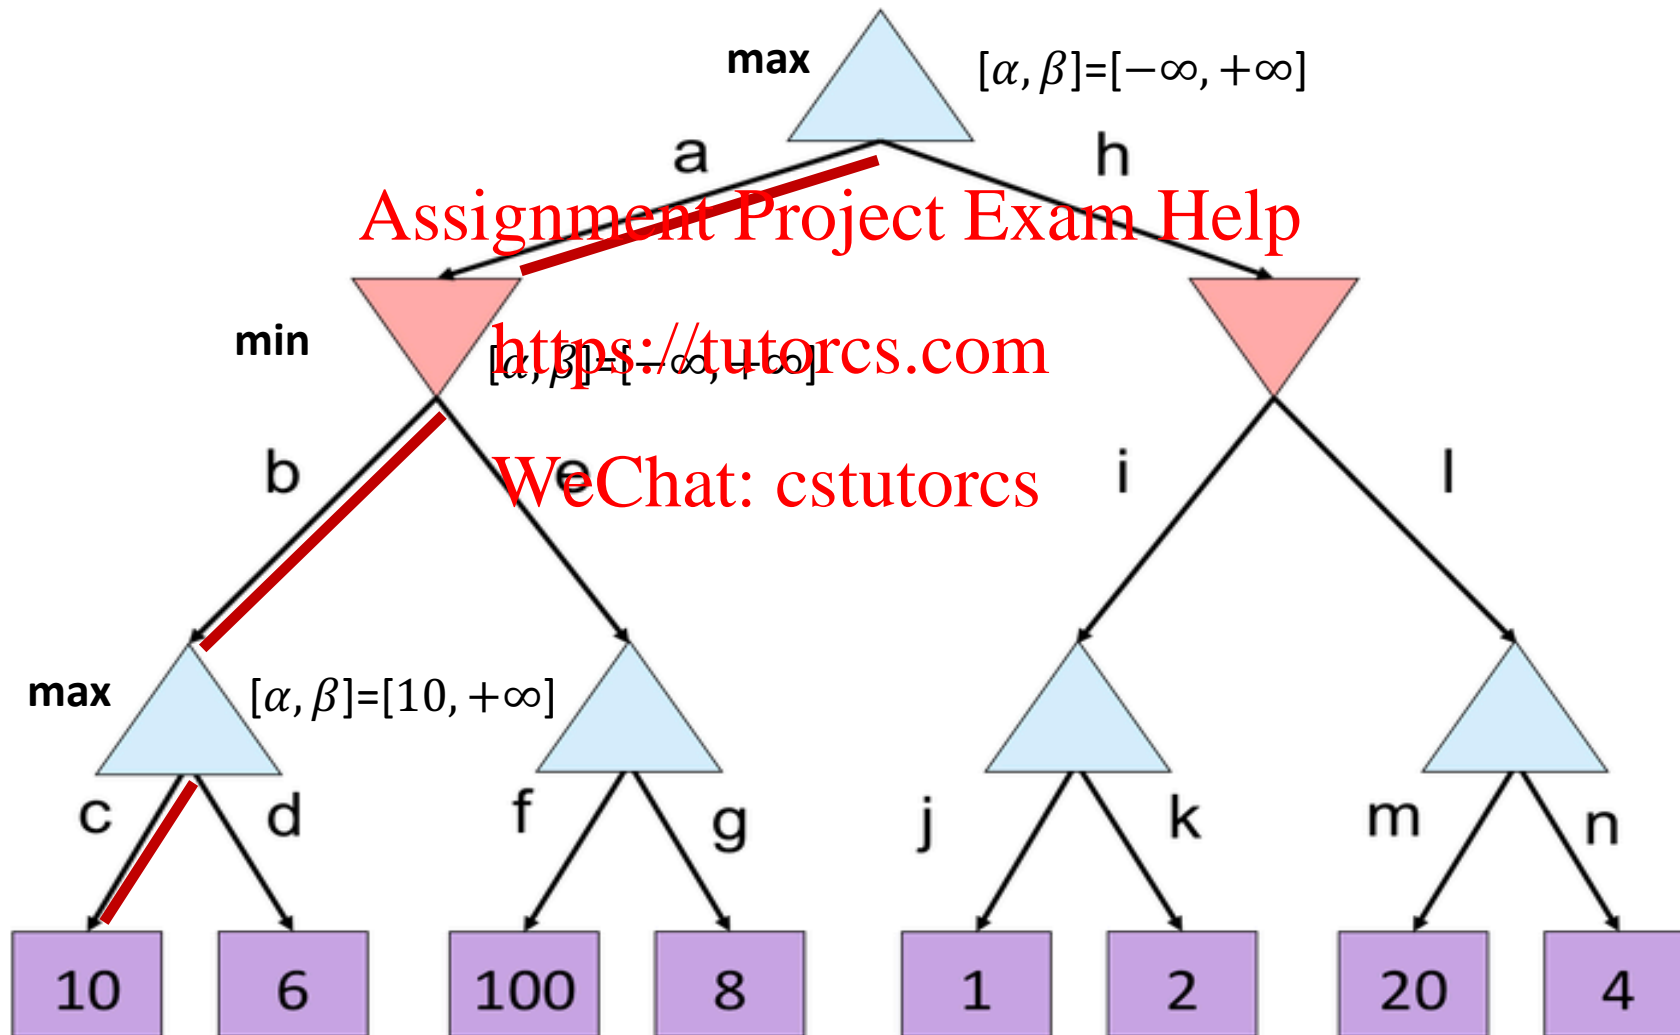
Alpha-Beta Quiz 2



Alpha-Beta Quiz 2



Alpha-Beta Quiz 2



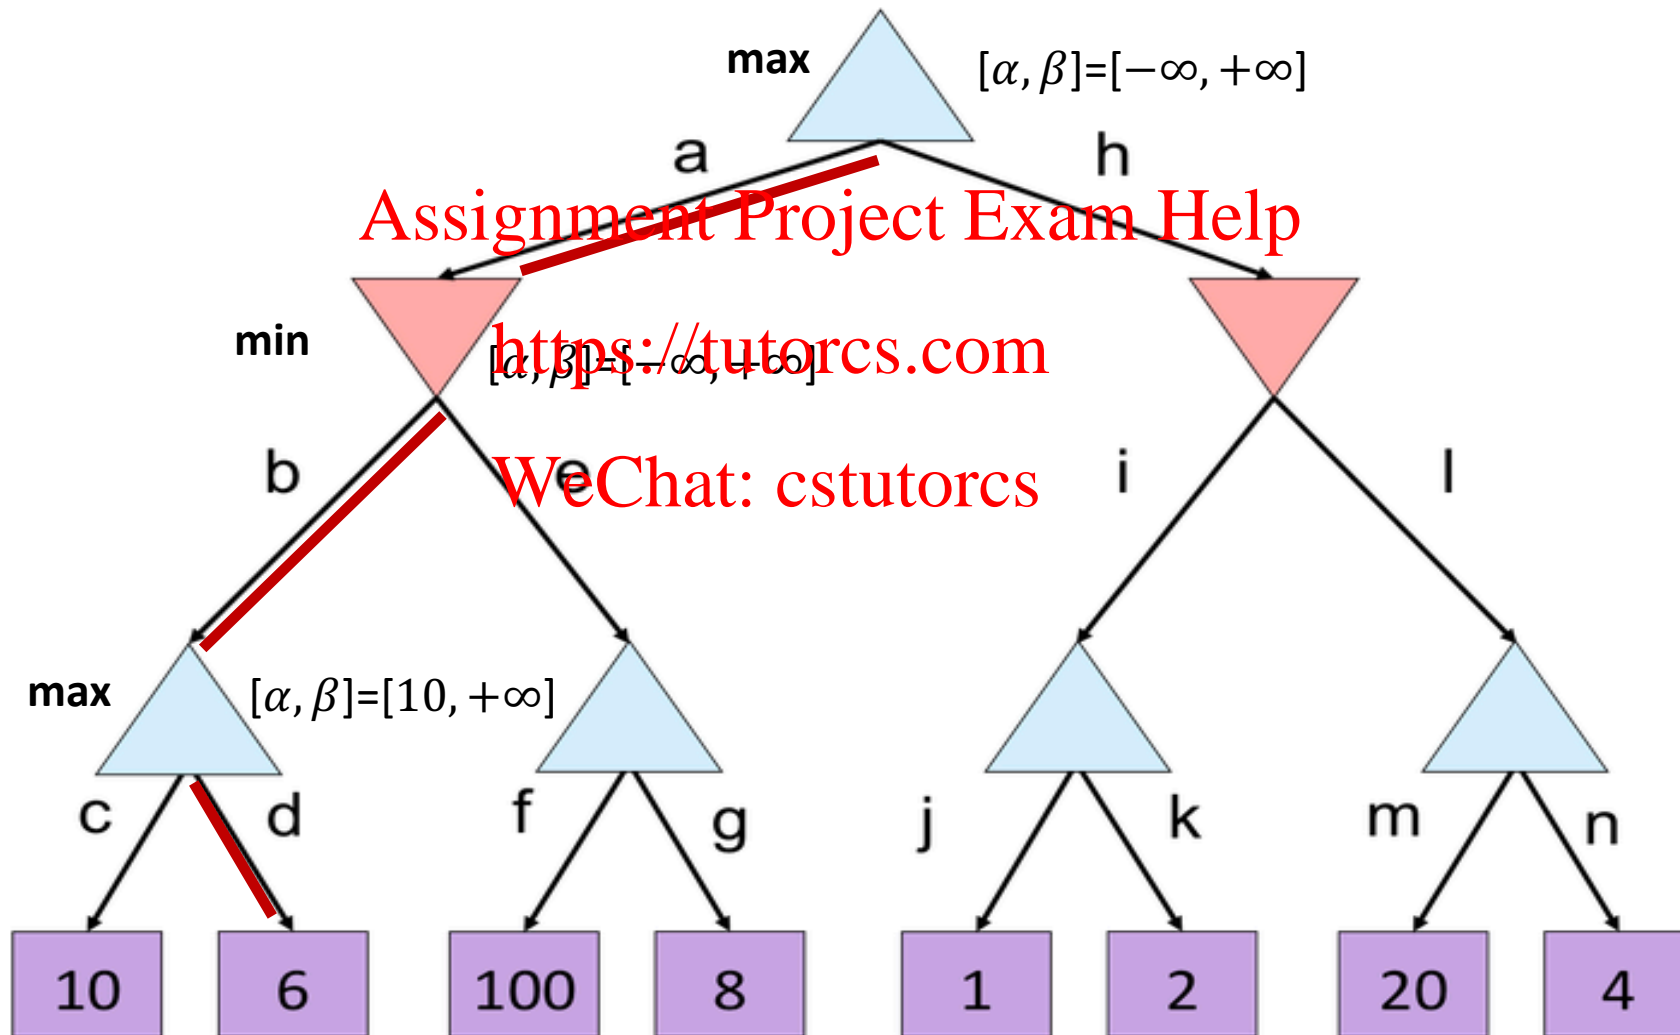
Assignment Project Exam Help

<https://tutorcs.com>

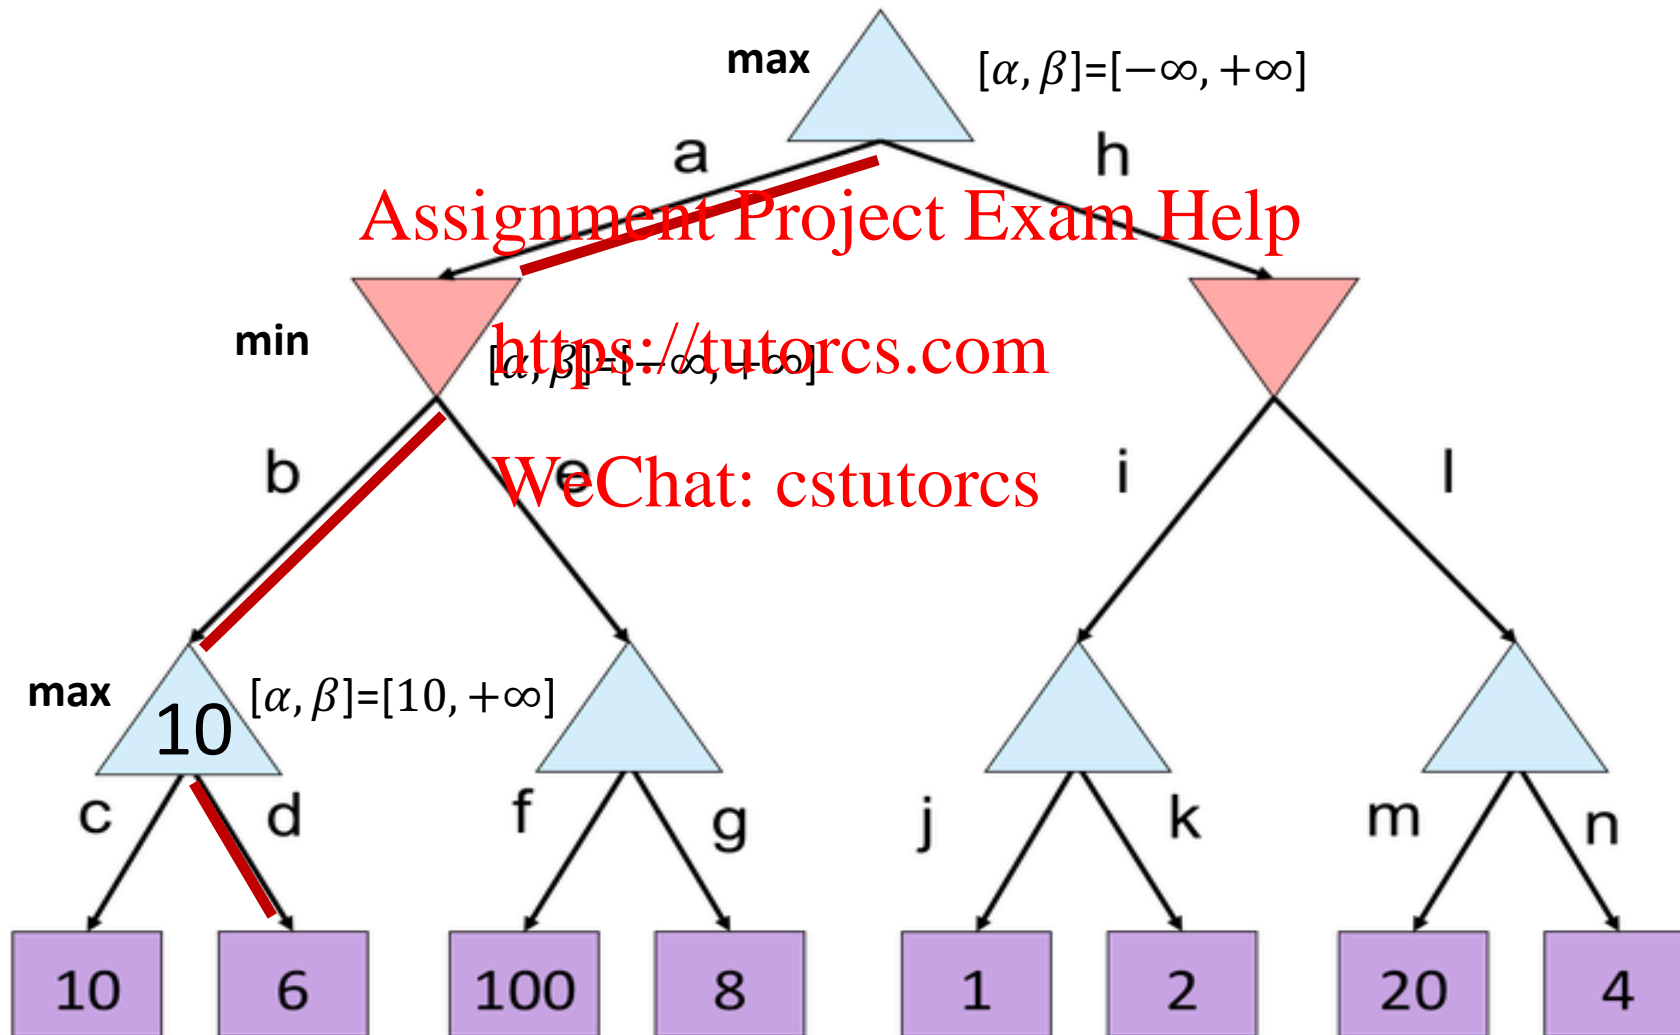
WeChat: cstutorcs



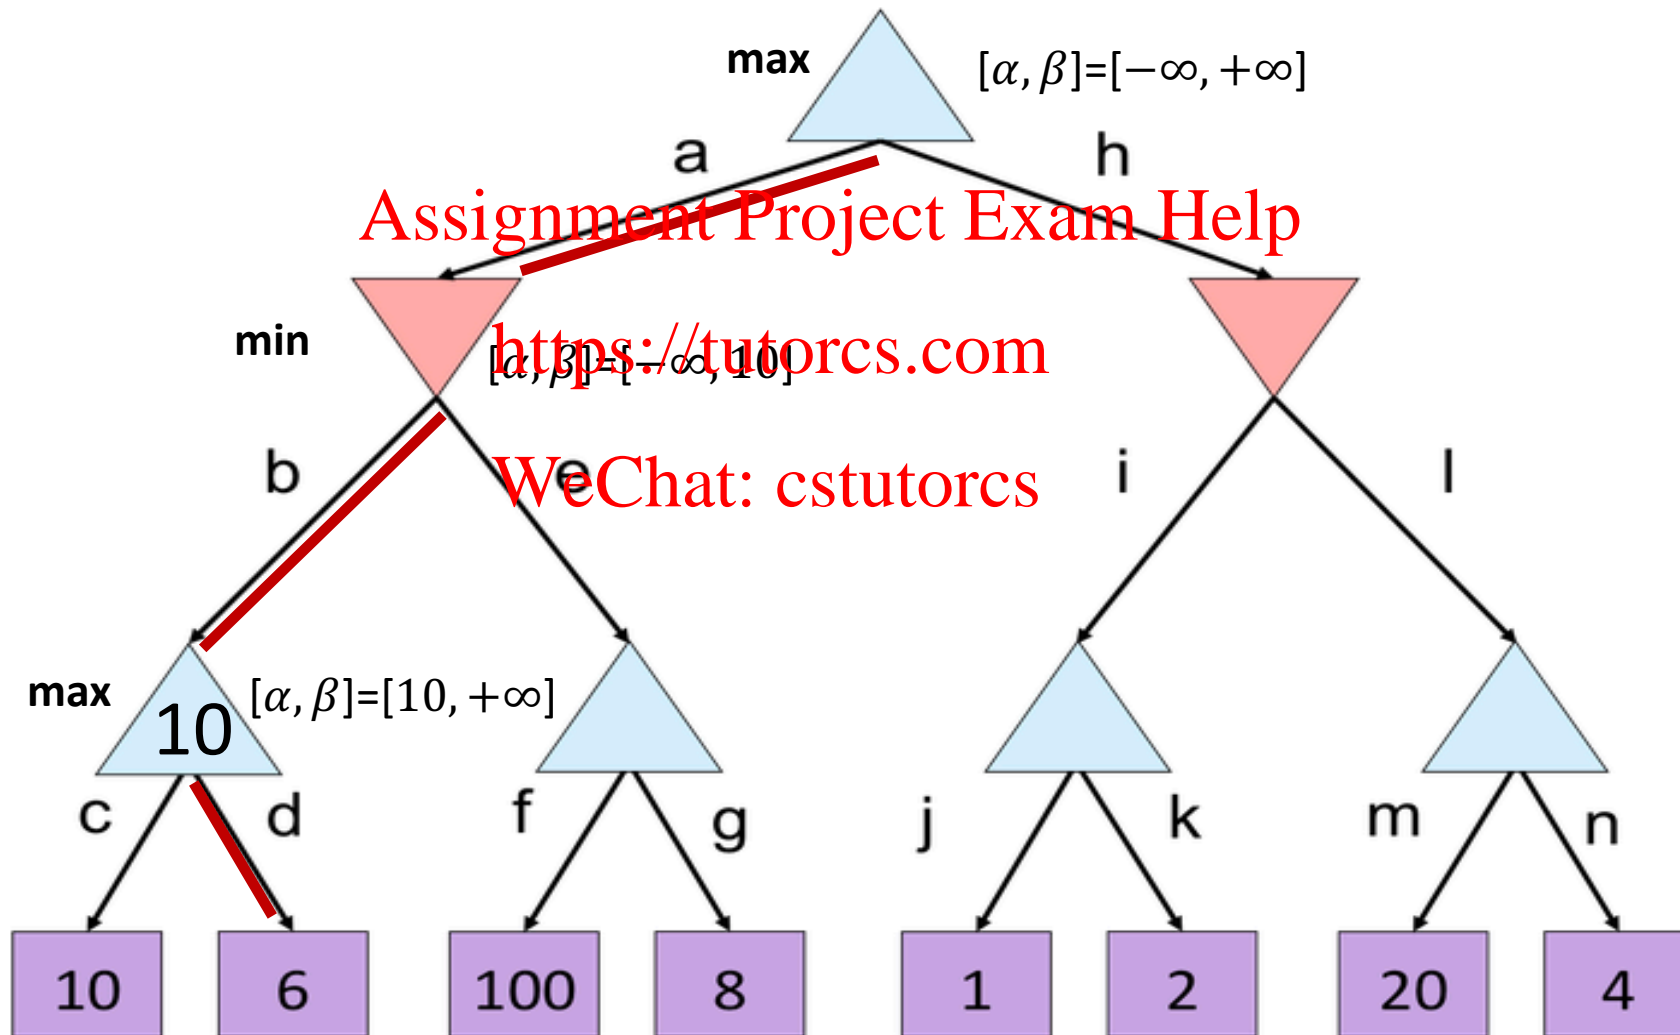
Alpha-Beta Quiz 2



Alpha-Beta Quiz 2



Alpha-Beta Quiz 2



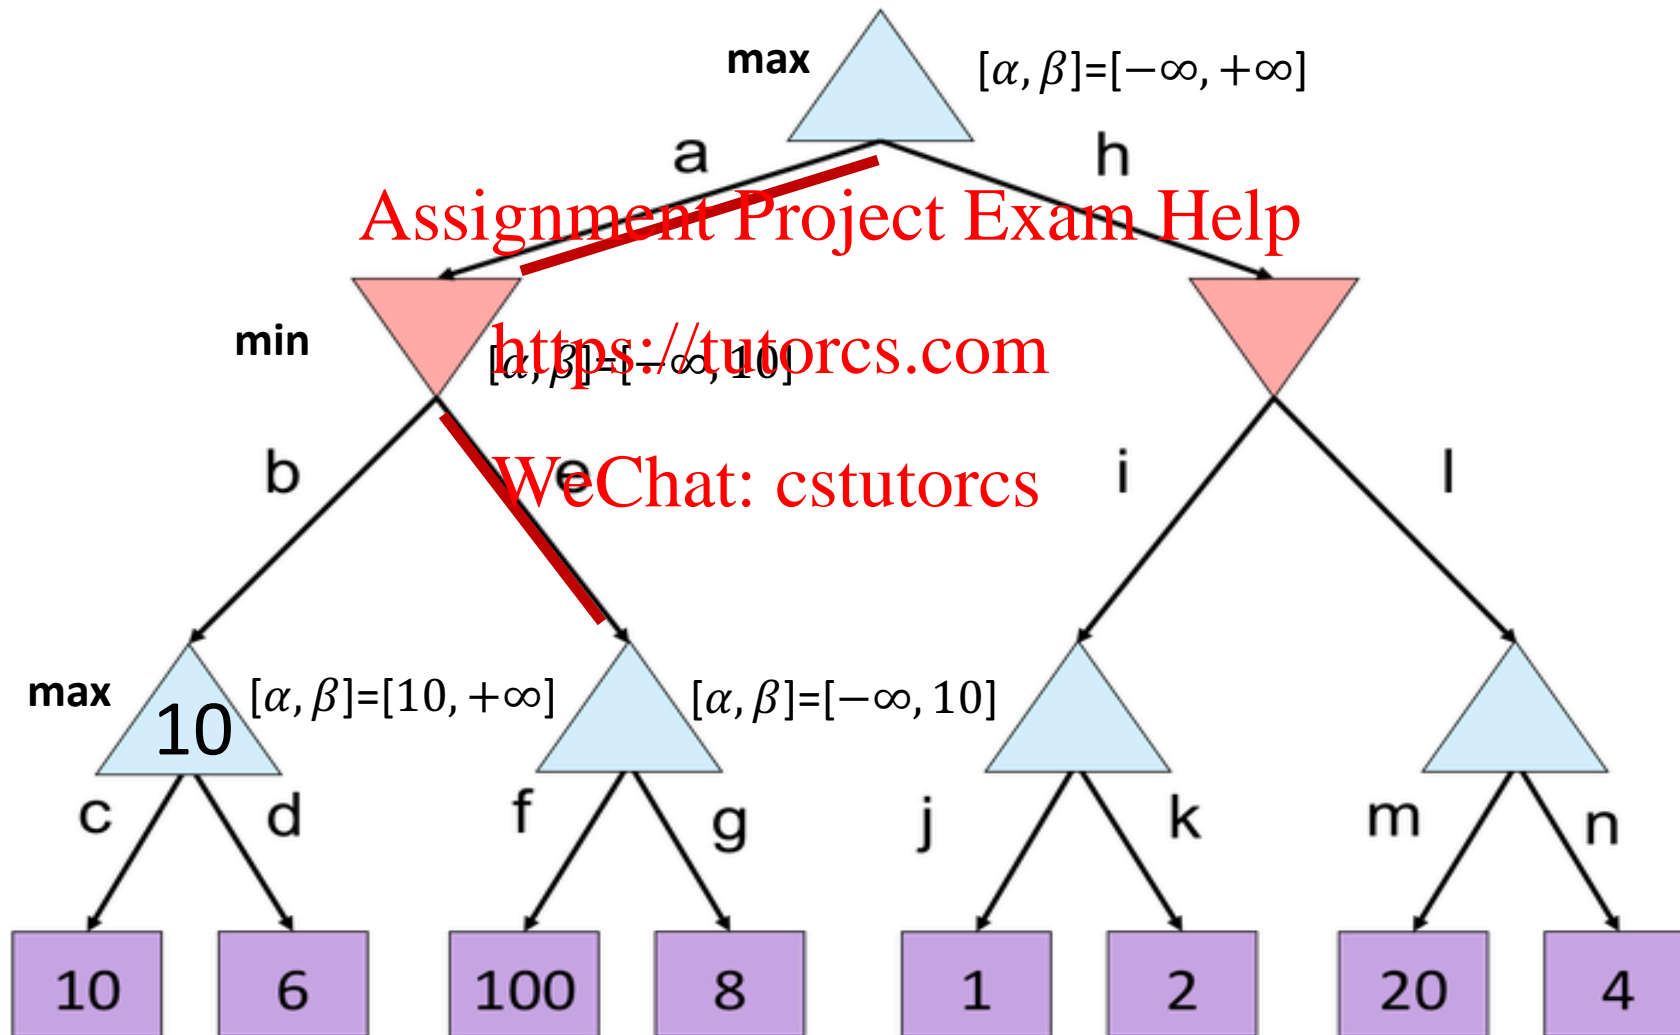
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Alpha-Beta Quiz 2



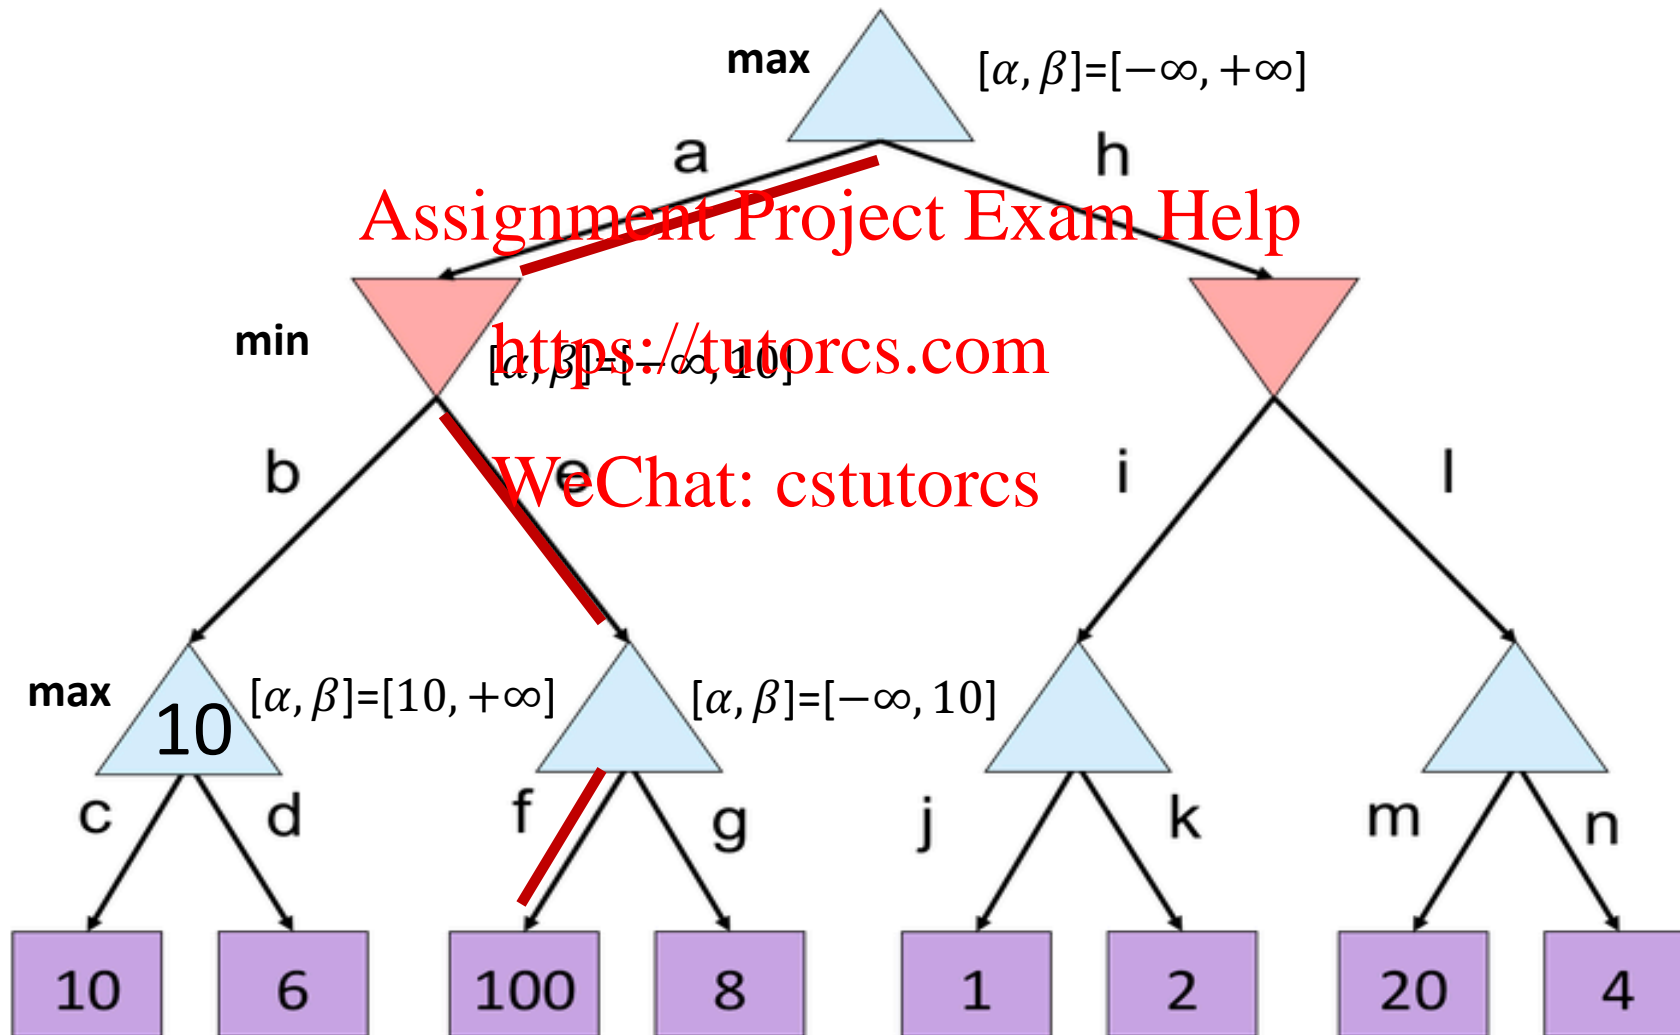
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Alpha-Beta Quiz 2



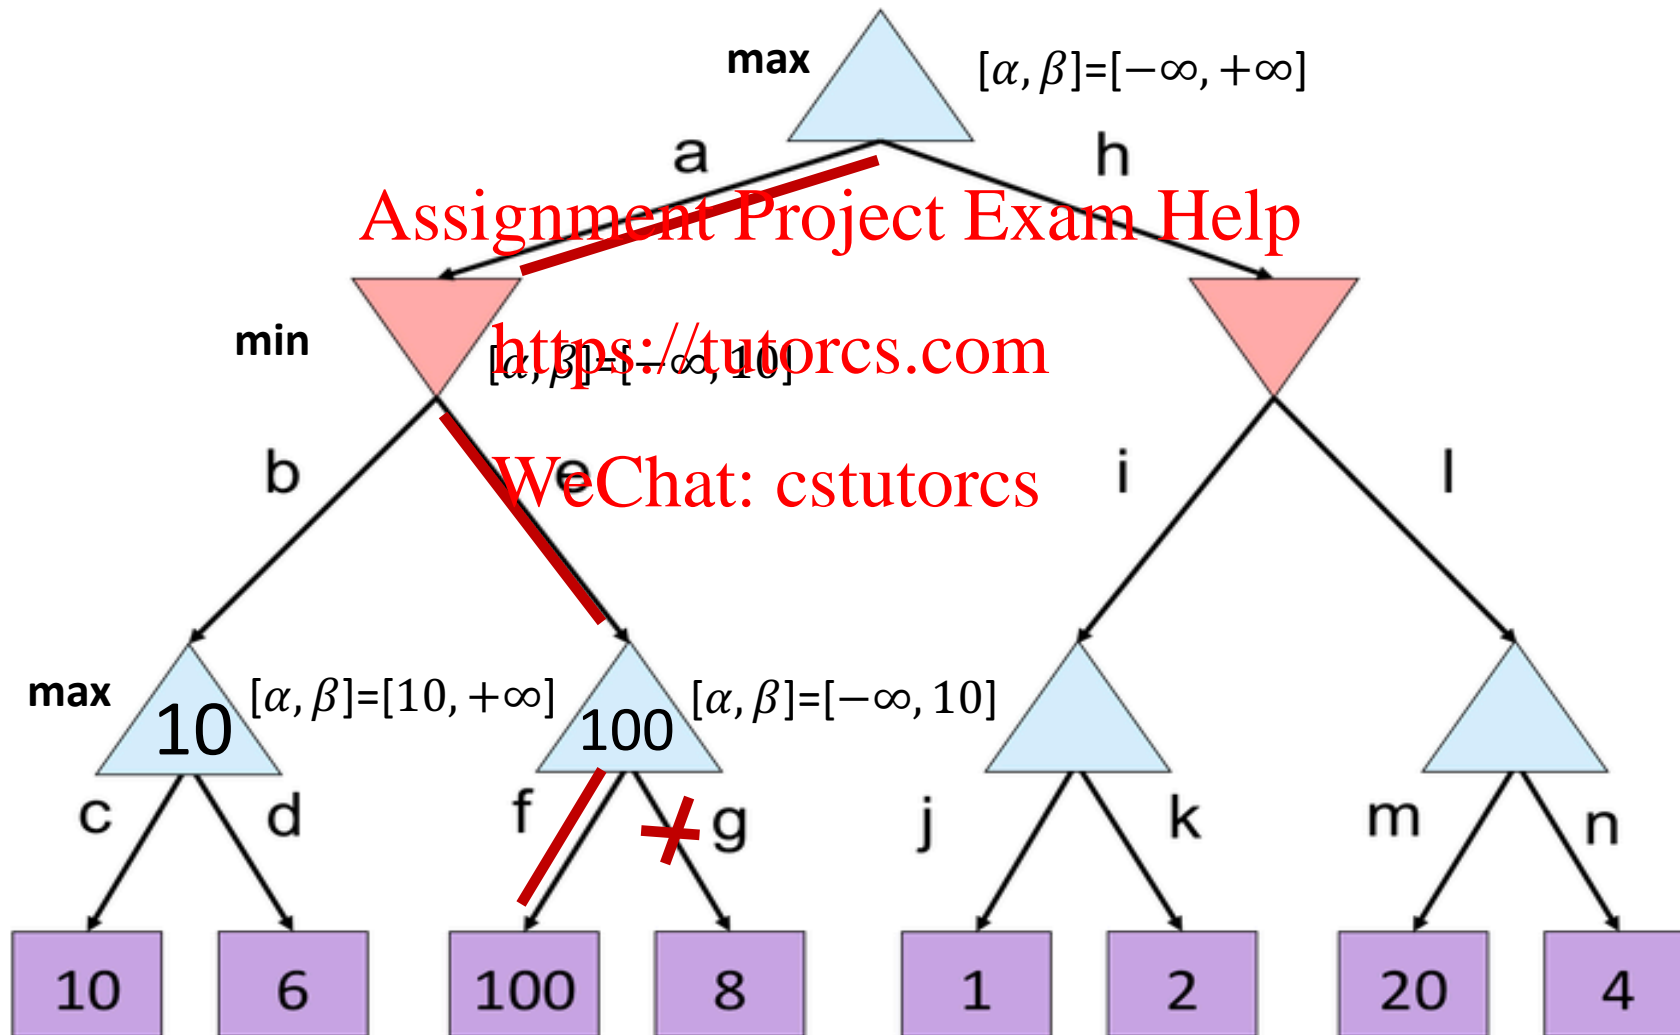
Assignment Project Exam Help

<https://tutorcs.com>

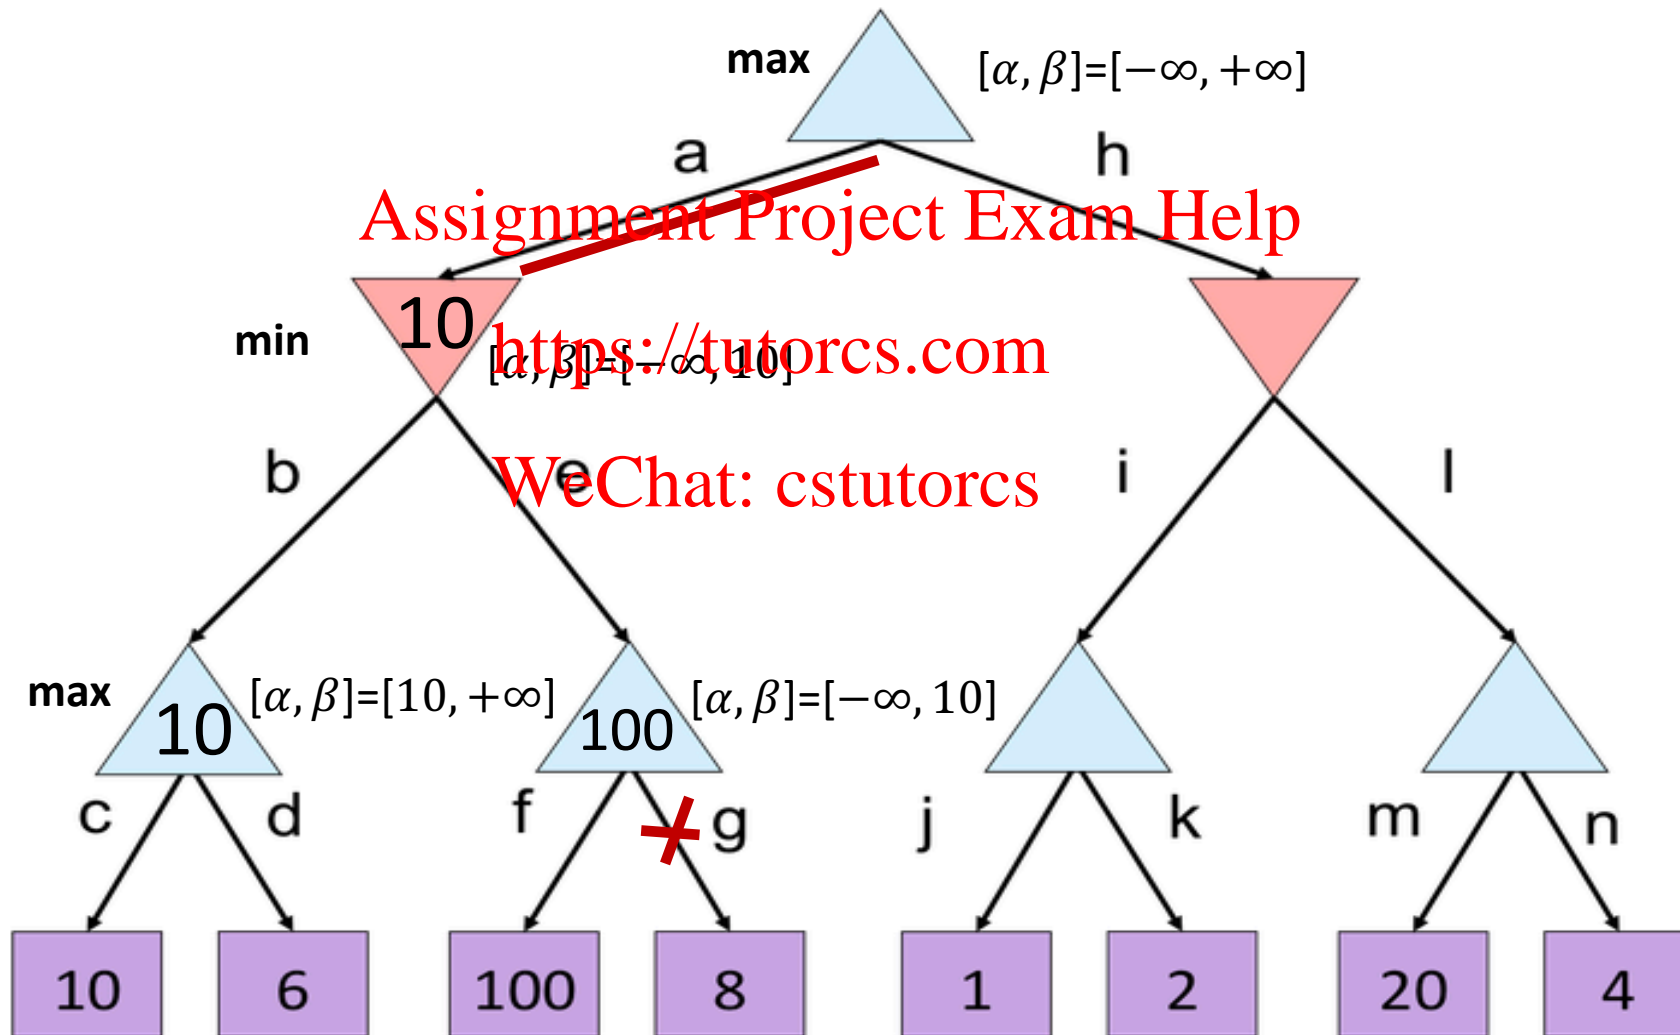
WeChat: cstutorcs



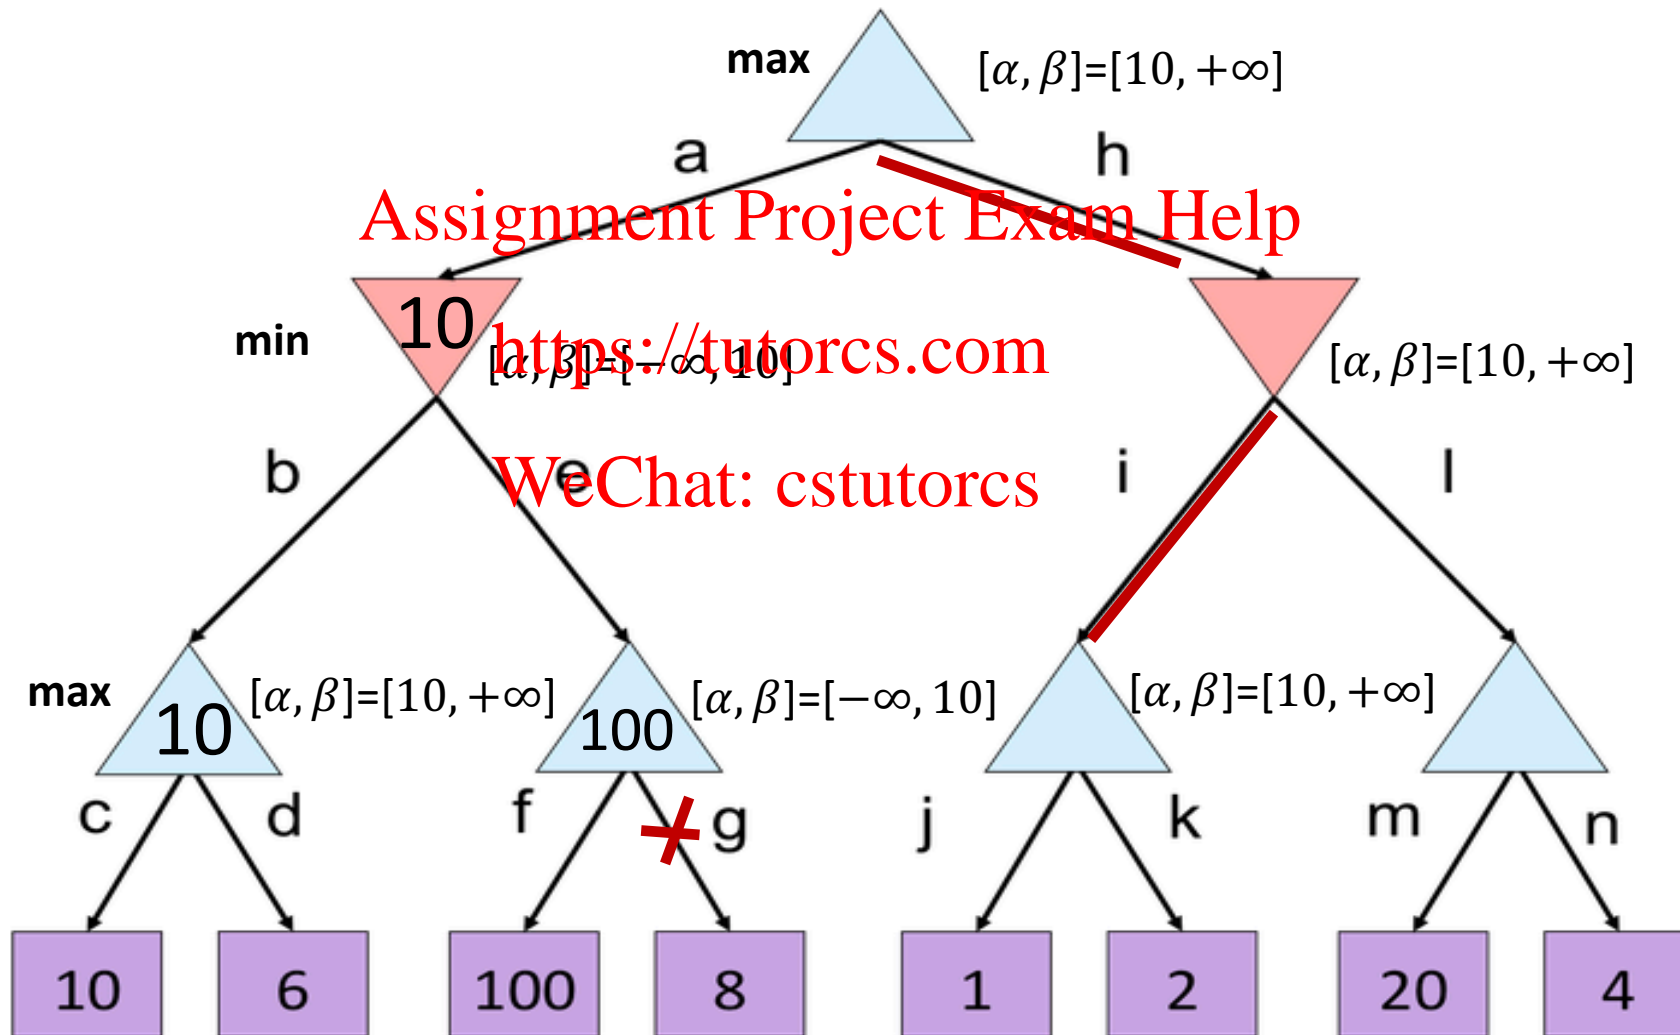
Alpha-Beta Quiz 2



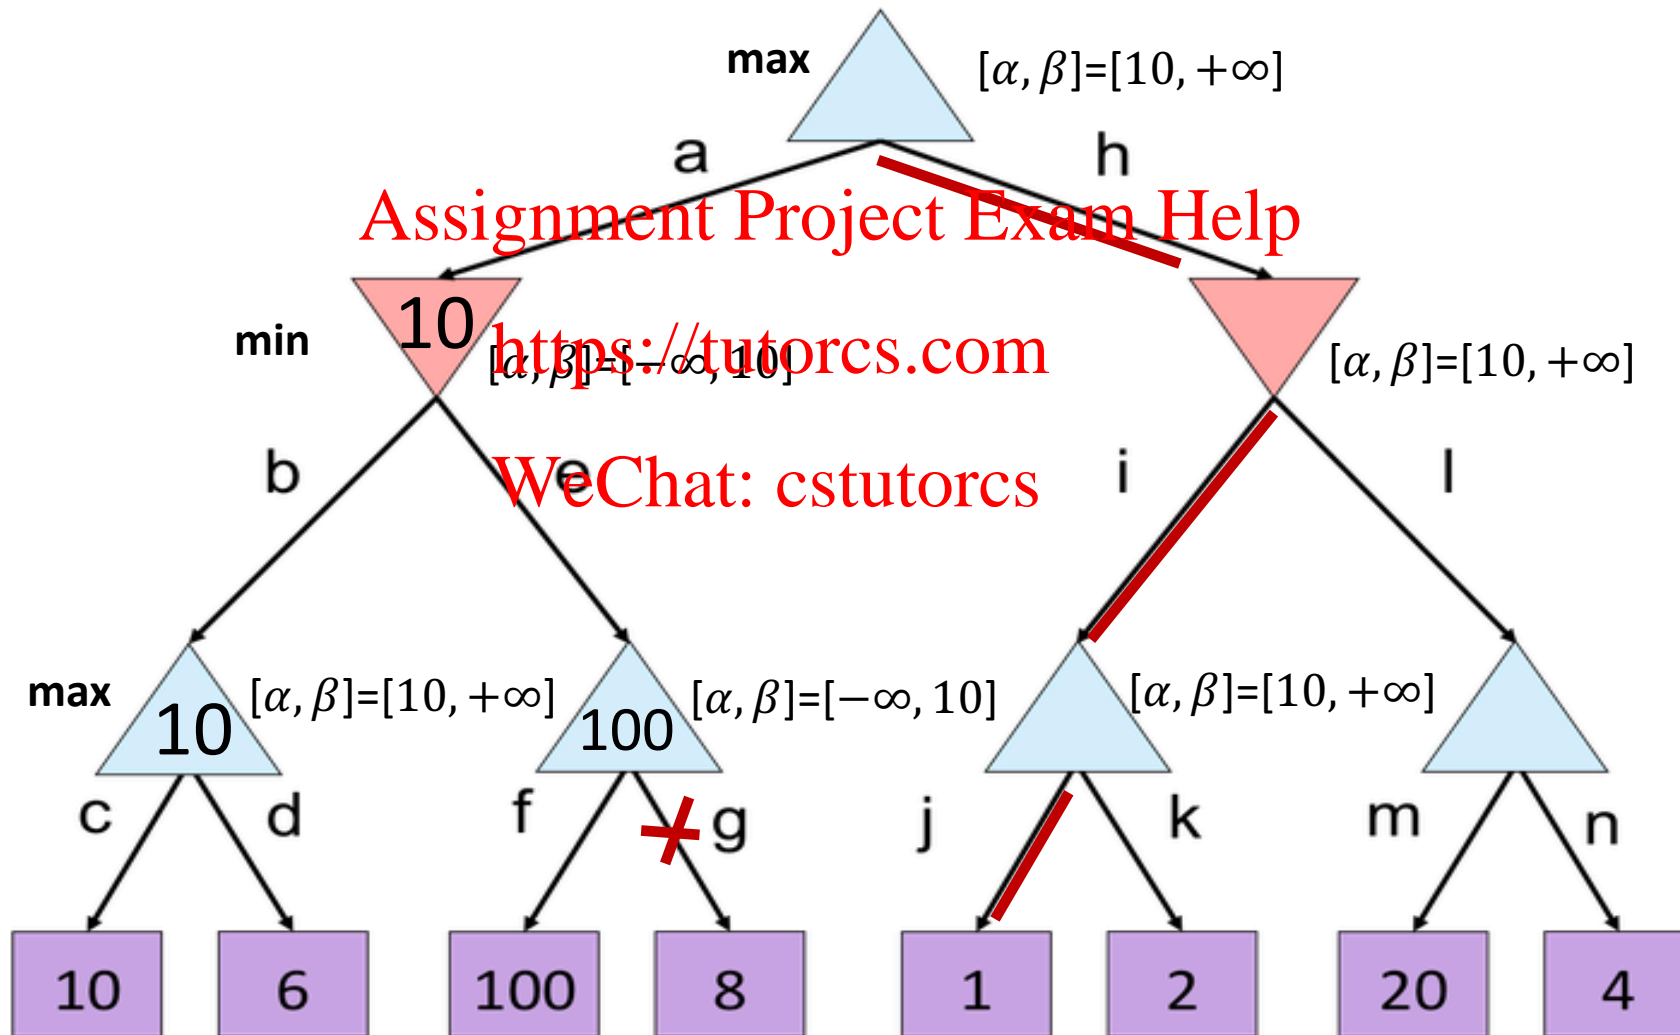
Alpha-Beta Quiz 2



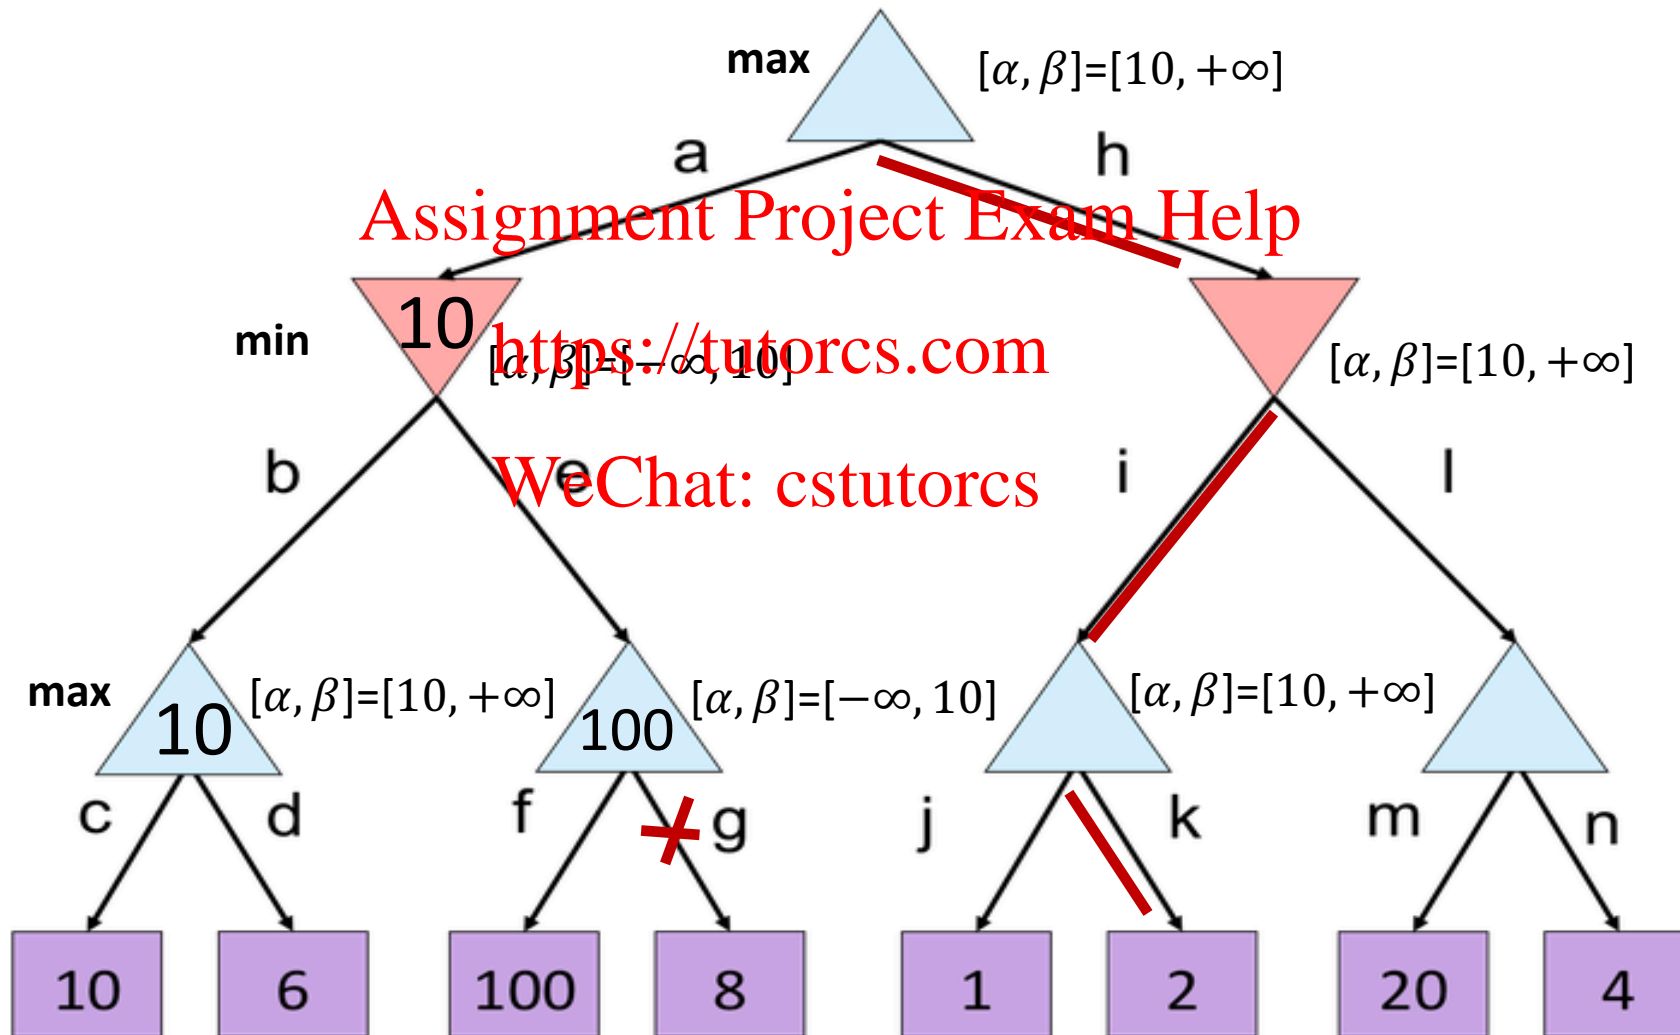
Alpha-Beta Quiz 2



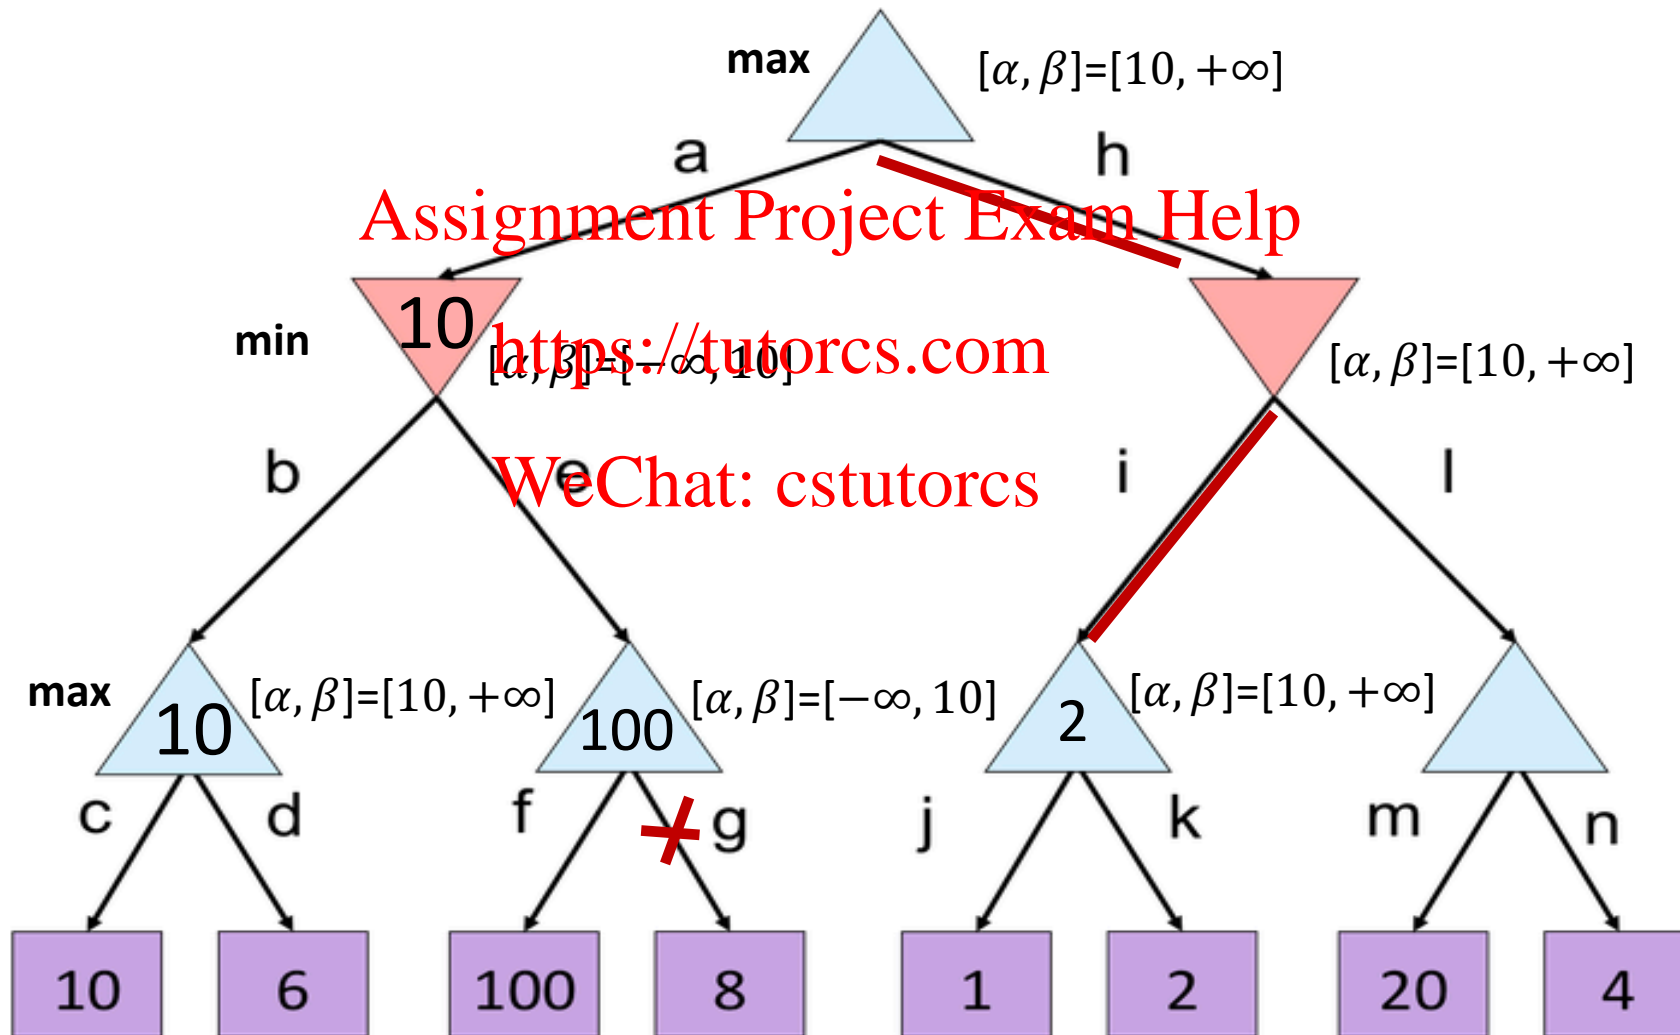
Alpha-Beta Quiz 2



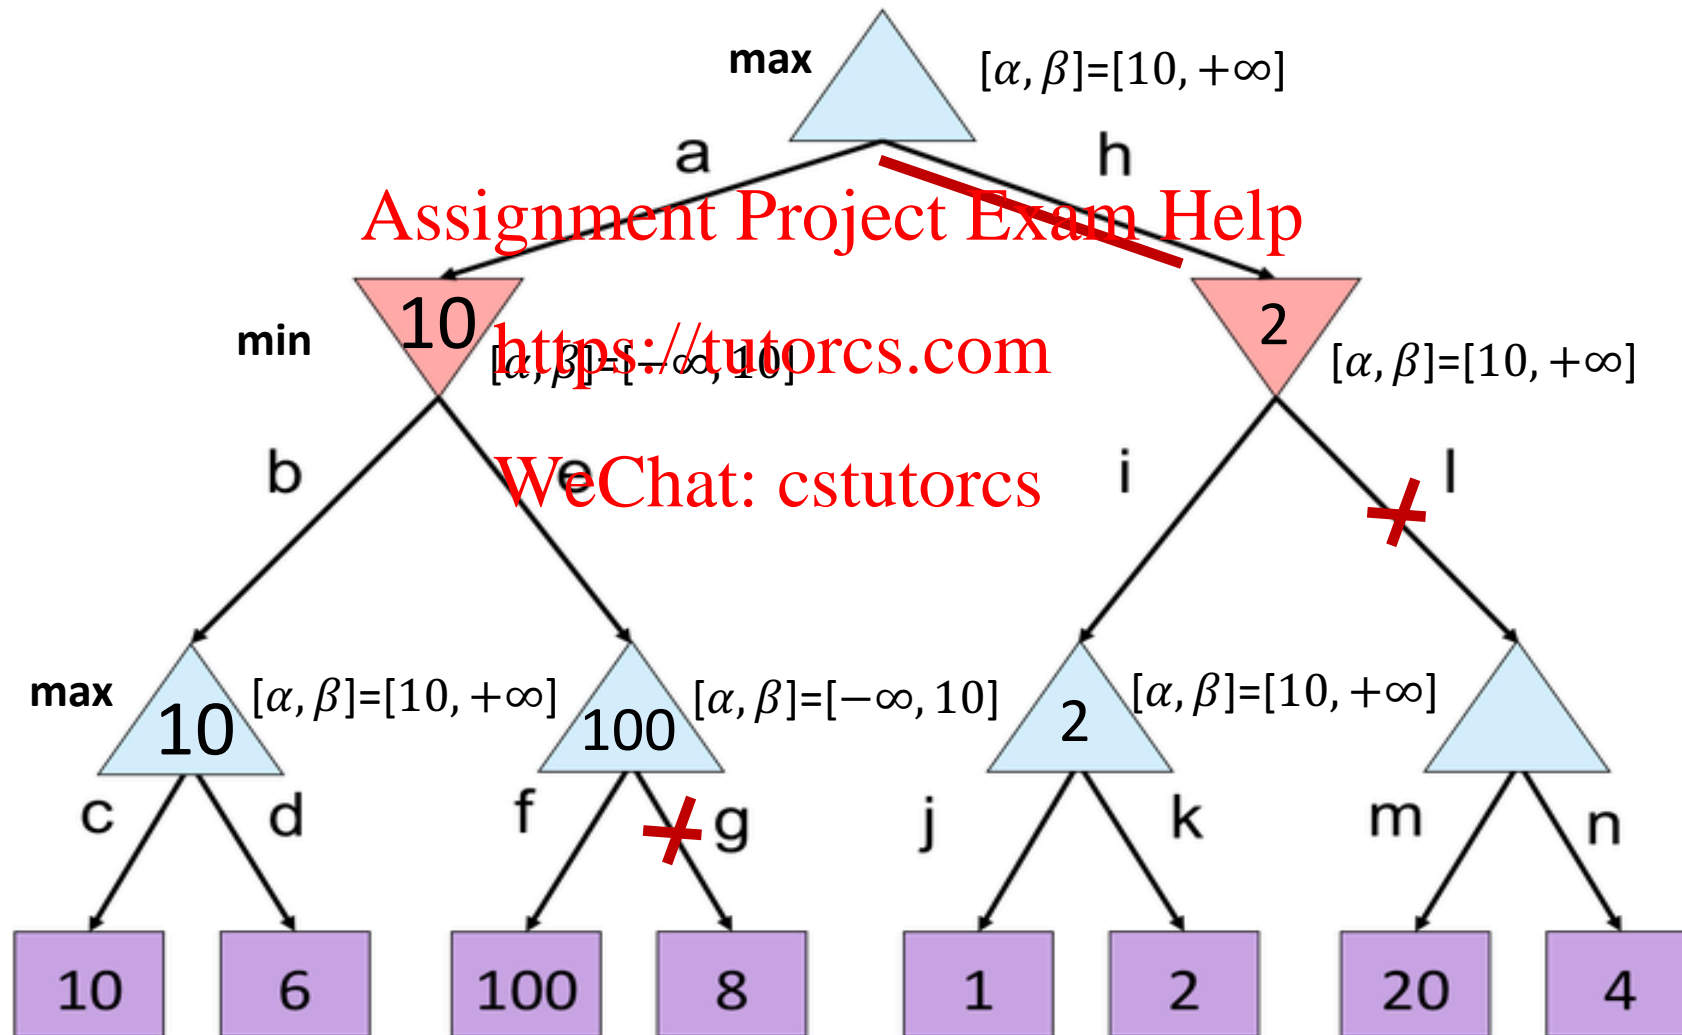
Alpha-Beta Quiz 2



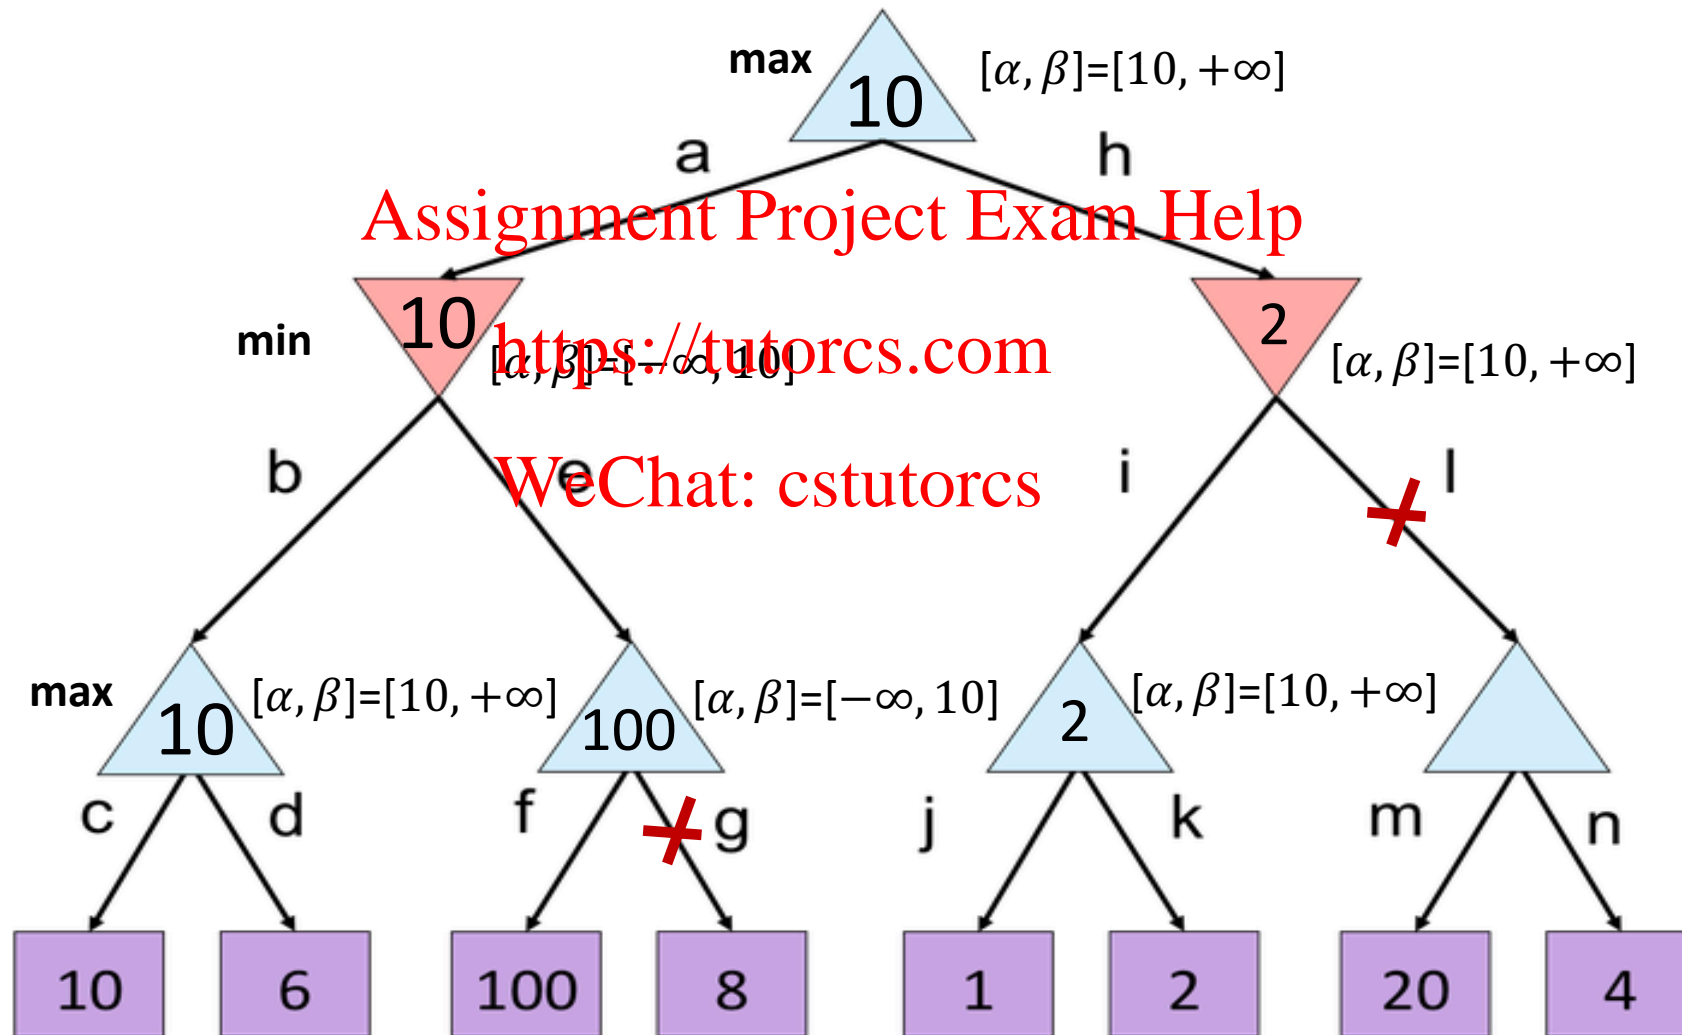
Alpha-Beta Quiz 2



Alpha-Beta Quiz 2



Alpha-Beta Quiz 2



Resource Limits

Assignment Project Exam Help


<https://tutorcs.com>

WeChat: cstutorcs



Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
 - Instead, search only to a limited depth in the tree
 - Replace terminal utilities with an evaluation function for non-terminal positions
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - α - β reaches about depth 8 – decent chess program
- Guarantee of optimal play is gone
- More plies makes a BIG difference
- Use iterative deepening for an anytime algorithm



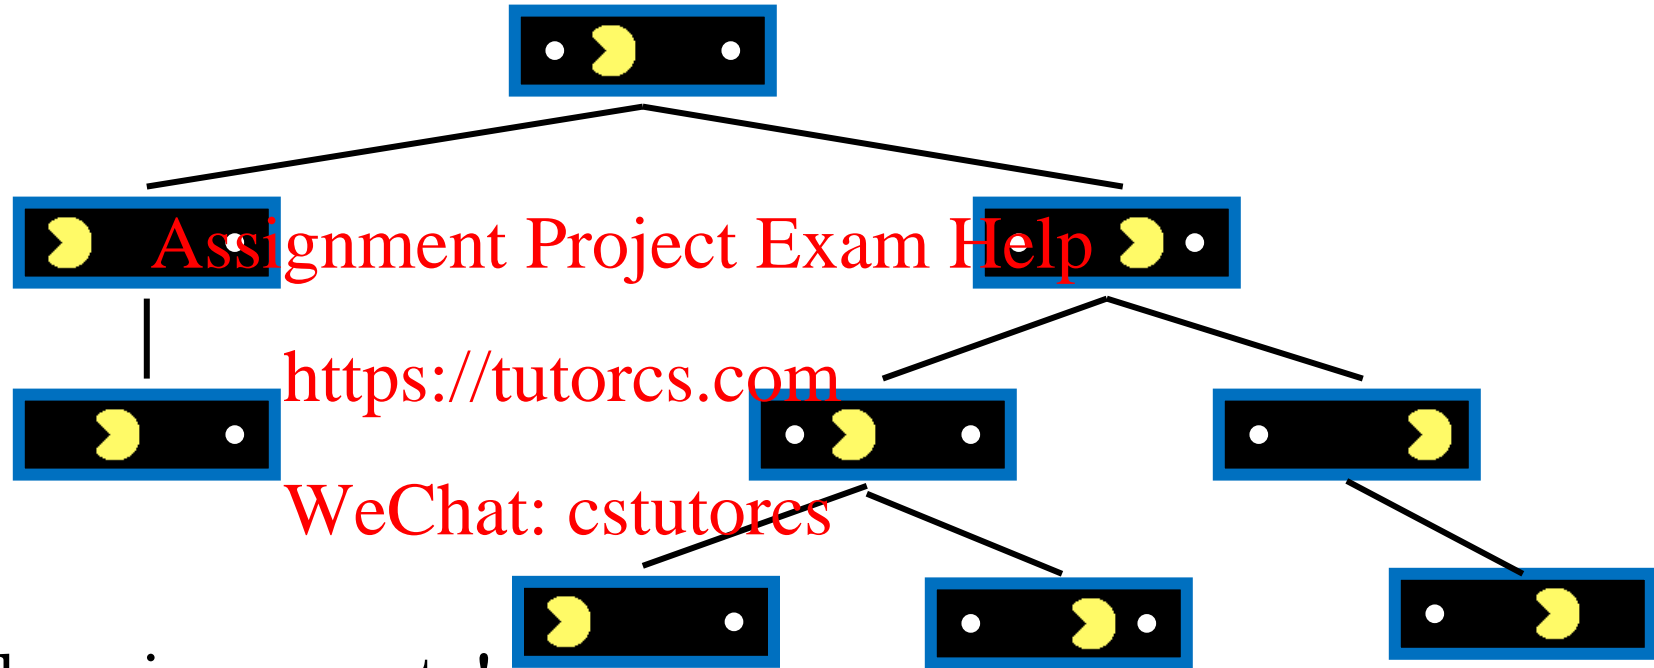
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



Why Pacman Starves



- A danger of replanning agents!
 - He knows his score will go up by eating the dot now (west, east)
 - He knows his score will go up just as much by eating the dot later (east, west)
 - There are no point-scoring opportunities after eating the dot (within the horizon, two here)
 - Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!




Evaluation Functions

Assignment Project Exam Help

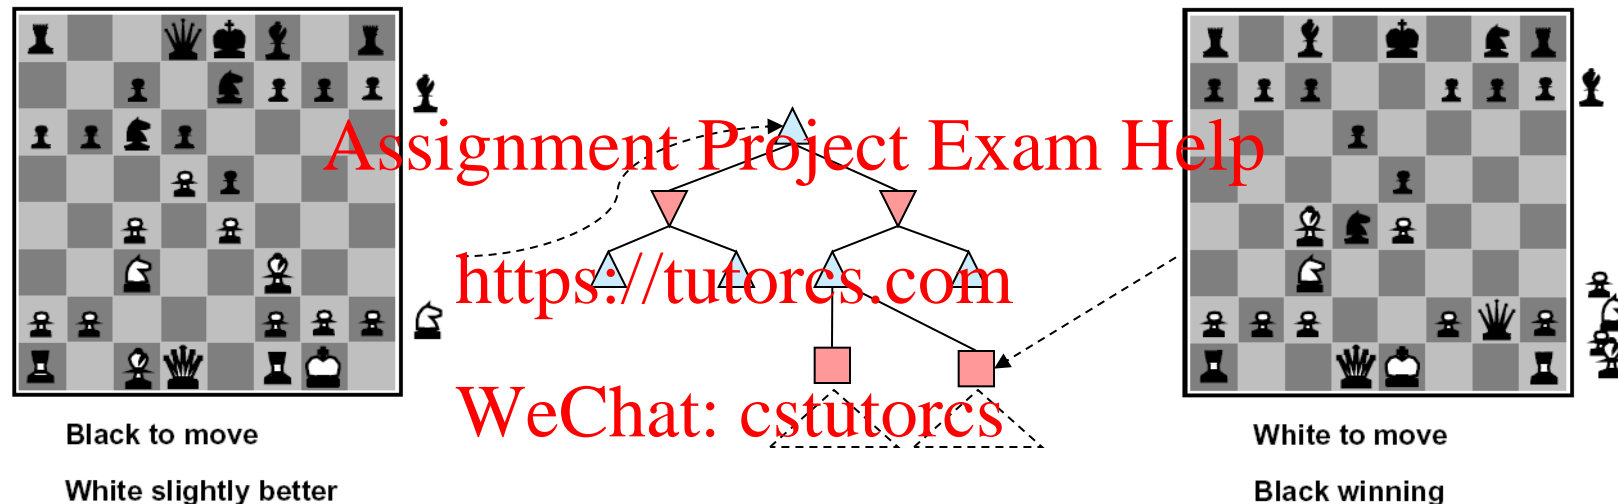
<https://tutorcs.com>

WeChat: cstutorcs



Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



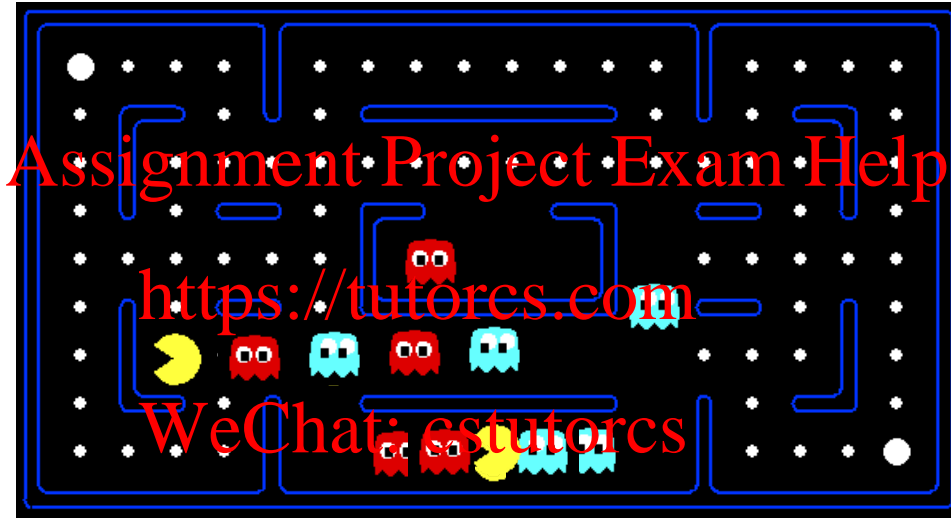
- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.




Evaluation for Pacman



Depth Matters

- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation



Synergies between Evaluation Function and Alpha-Beta?

- Alpha-Beta: amount of pruning depends on expansion ordering
 - Evaluation function can provide guidance to expand most promising nodes first (which later makes it more likely there is already a good alternative on the path to the root)
 - (somewhat similar to role of A* heuristic, CSPs filtering)
- Alpha-Beta: (similar for roles of min-max swapped)
 - Value at a min-node will only keep going down
 - Once value of min-node lower than better option for max along path to root, can prune
 - Hence: IF evaluation function provides upper-bound on value at min-node, and upper-bound already lower than better option for max along path to root THEN can prune

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

